

ArkPrism: Evidence-Carrying Privacy API Resolution for ArkTS Taint Analysis

ANONYMOUS AUTHOR(S)

Static privacy analysis for OpenHarmony/ArkTS depends on a step that conventional source–sink configuration leaves implicit: binding a platform API identity to the program value that carries its output. Package migration, manager receivers, property reads, and IR aliases can erase this identity before taint propagation begins; member-name matching can instead confuse unrelated application code with a platform source.

We present ARKPRISM, an evidence-preserving static analysis that resolves API identity from package, namespace, receiver, access-kind, and IR witnesses and binds the result to a typed carrier. ArkTS-specific IFDS recovery propagates these carriers through initializers, closures, exceptions, and access paths; a bounded supplementary analysis covers selected callback, Promise-then, and await continuations. Provenance-annotated evidence graphs retain the identity witness, call context, sink, path producer, and source location while keeping privacy-data and framework-input paths distinct.

On an output-independent source-first benchmark, ARKPRISM recovers all 90 gold occurrences at exact locations and rejects all 96 candidate collisions. It classifies all 64 matched identity cases correctly. On the complete 67-case HapBench suite, ARKPRISM achieves 100% case-level precision and specificity, 97.09% F1, and 98.04% endpoint-pair precision. A balanced asynchronous benchmark attributes nine Promise-then cases uniquely to supplementation. Manual source/IR review confirms all 90 sampled real-project privacy-data paths and 114/120 paths overall. Across 1,014 projects, the same frozen analysis identifies 2,336 privacy-API usages and 243 privacy-data paths without resource fallback.

CCS Concepts: • **Security and privacy** → **Software security engineering**; • **Software and its engineering** → *Software defect analysis*.

Additional Key Words and Phrases: OpenHarmony, ArkTS, static analysis, privacy APIs, taint analysis, call-chain analysis

ACM Reference Format:

Anonymous Author(s). 2027. ArkPrism: Evidence-Carrying Privacy API Resolution for ArkTS Taint Analysis. In *Proceedings of The ACM International Conference on the Foundations of Software Engineering (FSE 2027)*. ACM, New York, NY, USA, 34 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnn>

1 Introduction

OpenHarmony applications use ArkTS, a TypeScript-derived language with declarative ArkUI components, framework-managed callbacks, and system APIs exposed through historical `@ohos.*` and newer `@kit.*` packages [28, 32, 41]. These features complicate a task that precedes every privacy-flow query: identifying which program value is produced by a configured platform API. The same member token may denote application code, a third-party library, or a platform operation; the platform operation may appear as an imported namespace call, a property read, an IR alias, or a call on a manager returned by a factory.

ArkAnalyzer provides reusable IR, control-flow, call-graph, and pointer-analysis facilities for ArkTS [14], and HapFlow builds an IFDS taint analysis on this foundation [15]. Their availability makes whole-program ArkTS analysis practical, but it also exposes a front-end boundary: a solver can propagate only values that have first been bound to source identities. HapFlow’s released namespace–method loader maps configured entries to SDK method signatures. Its configuration schema does not encode executable property sources or retain the package/factory witness needed to decide a member on a manager receiver. Once source identity or its carrier is lost, downstream propagation cannot reconstruct the missing initial fact.

Consider `helper.createAsset(...)` after `helper` is returned by `photoAccessHelper.getPhotoAccessHelper(ctx)`. The member `createAsset` is common application vocabulary; its receiver origin establishes the platform identity. Conversely, accepting every same-named call in a file importing a media package would introduce collisions. Compound APIs such as `request.agent.create` create a related problem after IR lowering, while `deviceInfo.brand` produces a value through a property rather than a call. A useful analysis must therefore answer three linked questions: *which API is this*, *which value carries its output*, and *which evidence justifies both decisions*.

We present ARKPRISM, an evidence-preserving static analysis for ArkTS privacy auditing. For each candidate, it computes an API identity from package migration, namespace/import evidence, receiver type or origin, member, access kind, and IR aliases. Generic members such as `on`, `start`, and `close` require receiver or target evidence rather than lexical namespace proximity. The accepted identity determines a typed carrier—return value, property value, callback parameter, Promise resolution, or framework input—which seeds IFDS after ArkTS initializer, closure, exception, and access-path recovery. A bounded supplementary analysis recovers selected callback, Promise-then, and `await` continuations absent from core IR, while provenance records whether a path is produced by IFDS, supplementation, or both.

The resulting evidence graph preserves the acceptance witness, source carrier, call context, detector-local sink observations, configured-query paths, permissions, and source locations. It also keeps three evidence universes explicit: a recognized API usage need not produce a path; detector-local sinks are not the IFDS sink set; and framework-provided inputs are not silently counted as privacy-API outputs. This separation makes each reported edge auditable and prevents a broad lifecycle model from inflating the core privacy-data claim.

Our evaluation uses independent oracles for identity, flow classification, mechanism attribution, and scale. On a source-first real-project benchmark, ARKPRISM recovers all 90 gold occurrences at exact locations and rejects all 96 candidate collisions. It classifies all 64 matched identity cases correctly. On the complete 67-case HapBench suite, ARKPRISM achieves 100% case-level precision and specificity, 97.09% F1, and 98.04% endpoint-pair precision. ArkAsyncBench attributes nine Promise-then cases uniquely to asynchronous supplementation. A source/IR audit of 120 stratified real-project paths confirms all 90 sampled privacy-data paths and 114/120 paths overall; the six rejected candidates are confined to the separately labeled framework-input layer. Across 1,014 projects, the frozen analysis identifies 2,336 privacy-API usages and 243 privacy-data paths, with no resource fallback.

This paper makes four contributions:

- **Evidence-carrying API identity.** We formulate ArkTS privacy-API recognition over package migration, receiver evidence, access kind, and IR aliases, covering direct calls, manager/helper receivers, compound namespaces, and executable properties while retaining an explicit match witness.
- **Carrier-aware propagation.** We bind each identity to a typed source carrier, recover ArkTS-specific IR semantics inside IFDS, and supplement unresolved asynchronous continuations under explicit bounds and provenance.
- **Auditable evidence graphs.** We assemble identity, call context, sinks, paths, metadata, and source locations without conflating detector evidence, privacy-data flows, and framework-input flows.
- **A multi-level empirical evaluation.** We contribute source-first identity benchmarks, a complete HapBench reproduction, a balanced asynchronous benchmark, a 120-path semantic

audit, and a clone-aware 1,014-project study whose tables and figures are tied to machine-readable manifests.

2 Background and Motivation

2.1 OpenHarmony and ArkTS

OpenHarmony applications are organized around abilities, pages, ArkUI components, lifecycle methods, and event callbacks [27, 28, 41]. ArkTS preserves much of the TypeScript surface syntax, but introduces stricter typing and a declarative UI model. In ArkUI, a component’s build method describes UI structure, while lifecycle callbacks such as `aboutToAppear` and event handlers such as `onClick` are implicitly invoked by the framework. These properties make ArkTS attractive for application development, but complicate static analysis because important control-flow edges are not always explicit in source code.

Unlike traditional Android applications, where many lifecycle and callback conventions have been studied for years, OpenHarmony combines a relatively new application model with a rapidly evolving SDK. An ArkTS page may contain ordinary methods, lifecycle methods, anonymous functions, class-field callbacks, Promise handlers, and declarative UI builder code in the same file. The compiler and analyzer may lower these constructs into generated methods such as `%AM-style` anonymous methods, `%df1t`, `%statInit`, and `%instInit`. These generated methods are not noise for privacy analysis: they often contain callback assignments, page initialization, or API wrapper setup. An analysis that filters them out too aggressively can lose real paths.

OpenHarmony also exposes a broad API ecosystem for device identity, account information, clipboard access, sensors, camera, media, location, network state, Wi-Fi, Bluetooth, telephony, calendar, and other capabilities. Some APIs have migrated from historical `@ohos.*` packages to newer `@kit.*` packages, while many apps still contain mixed imports, aliases, default imports, and manager-style APIs. For example, a project may import device information from `@ohos.deviceInfo`, from `@kit.BasicServicesKit`, or through a default alias. Similarly, APIs such as `pasteboard` and `request-agent` operations may be invoked directly from a namespace or indirectly through an object returned by a manager factory.

This evolution creates a tension: exact package matching separates historical and consolidated exports of the same API, while method-name-only matching accepts common verbs without an ownership witness. A resolver must normalize documented migrations at namespace granularity while preserving receiver and member constraints.

2.2 Foundational Static Analysis for ArkTS

ArkAnalyzer provides ArkTS-aware code representation, IR transformation, Scene construction, call-graph construction, and basic data-flow support [14]. HapFlow builds on ArkAnalyzer to provide inter-procedural taint analysis for OpenHarmony, addressing lifecycle/callback entry modeling, closure-captured variables, and source/sink extraction [15].

ArkAnalyzer and HapFlow play different roles: ArkAnalyzer is a representation and graph-construction layer (Ark files, classes, methods, CFG, imports, call graphs); HapFlow is a taint-analysis layer (whether configured sources reach configured sinks). ARKPRISM is built between and around these layers, asking a different question: can we recover the evidence that an auditor needs to understand a privacy-sensitive behavior?

Privacy auditing needs a *behavior-level* explanation: what API is used, how it is reached, what exposure points are downstream, what category and permission are involved, and how patterns

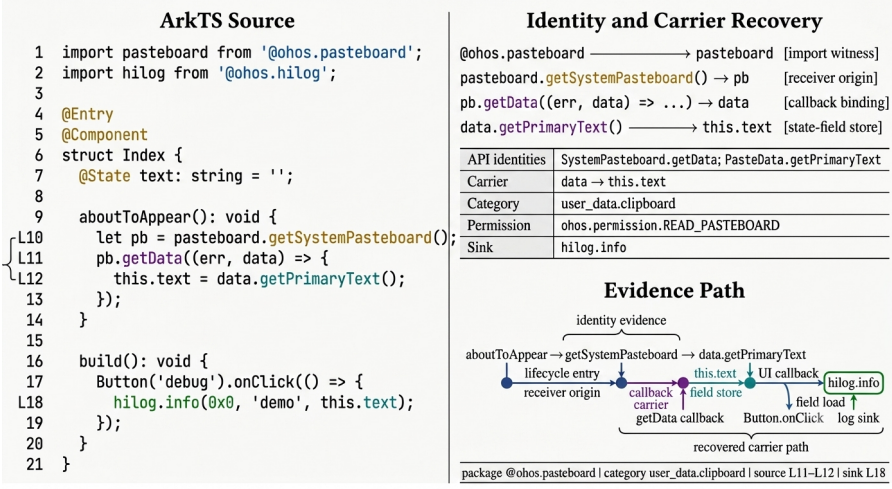


Fig. 1. Motivating ArkTS example and the evidence recovered by ARKPRISM.

appear across a corpus. This distinction affects evaluation: a taint benchmark tests crafted source-to-sink cases, while a corpus audit needs API recognition against source evidence, call-chain coverage, sink distribution, and manual confirmation.

2.3 Information-flow Subgraphs

For each recognized privacy-sensitive API usage u , ARKPRISM constructs an information-flow subgraph G_u that connects four evidence layers: the API usage itself, the calling context that reaches it, downstream exposure points (sinks), and semantic metadata (package, permission, category). The formal definition is in Section 3. This evidence structure deliberately separates several levels: a privacy-sensitive API usage may be real even when no downstream sink is found; a call chain may be locally precise even when no trusted entry can be recovered; a taint path may indicate data reachability but not necessarily a policy violation.

2.4 Motivating Example

Figure 1 shows a simplified ArkTS pattern. The app obtains the system pasteboard, reads data in a callback, stores the text in a component state field, and later logs it from a UI interaction. A source-to-sink analyzer should report the data flow. An information-flow subgraph should additionally explain that the source belongs to the clipboard category, that it is reached from a component lifecycle or callback context, that the sink is log output, and that the source is associated with a package/namespace import.

This example illustrates why a single analysis layer is insufficient: recognition must handle the indirect receiver pb, call-chain recovery must include lifecycle and UI callback edges, taint analysis must handle callback-provided source data, and sink analysis must classify hilog.info as log exposure. ARKPRISM combines these layers in one report.

2.5 Analysis Requirements

The ArkTS setting induces six recurring analysis requirements: robust project-layout discovery, namespace-constrained package migration, compound-namespace recovery after IR lowering, manager-receiver provenance, callback-edge recovery, and typed propagation of callback-carried

values. Section 3 maps these requirements to explicit analysis stages, while Section 5 tests them with identity, flow, and asynchronous oracles.

3 Design of ARKPRISM

ARKPRISM resolves the *source identity* bottleneck in ArkTS privacy analysis and produces provenance-annotated evidence subgraphs for auditing. ArkAnalyzer provides the ArkTS IR, project Scene, and call-graph substrate [14]; HapFlow provides IFDS-based taint propagation over OpenHarmony programs [15]. On this substrate, ARKPRISM adds an evidence-carrying domain layer: package- and receiver-aware API identity resolution, bounded asynchronous carrier/context recovery, and provenance-preserving subgraph assembly.

3.1 Problem Definition and Notation

Let a project be $P = \langle F, M, I, E_c \rangle$, where F is the set of ArkTS source files, M is the set of methods in ArkAnalyzer IR, I is the set of import bindings, and $E_c \subseteq M \times M$ is the initial call graph. ARKPRISM deliberately separates detector rules R_D from IFDS query rules R_T . A detector rule is

$$r_D = \langle pkg, ns, mem, direct, factory, perm, cat \rangle, \quad (1)$$

where pkg and ns identify the declaring package and namespace, mem is a method or property token, $direct \in \{true, false, null\}$ distinguishes namespace calls, receiver calls, and property reads, $factory$ is an optional set of receiver-producing methods, $perm$ is permission evidence, and cat is a privacy category. An IFDS query rule is

$$r_T = \langle sig, meta, carrier, i_{cb}, i_{src}, kind \rangle, \quad (2)$$

where sig is an SDK method signature; $meta = \langle module, ns, class, mem, params, lit \rangle$ retains the source identity needed when the frontend cannot resolve that signature; and $carrier \in \{return, arg, callback\}$ identifies the program value to seed. The source kind distinguishes typed privacy data from explicit framework inputs:

$$kind \in \{privacy_data, framework_input\}. \quad (3)$$

The optional lit component records a string-dispatched event name. For callback sources, i_{cb} is the callback-argument position and i_{src} is the sensitive callback-parameter position. A privacy-data return requires a non-void return type; a callback source requires a typed `Callback<T>` parameter and a valid carrier index; an argument source is accepted only as an explicitly modeled framework input. Rules without this carrier contract are rejected at load time. Keeping R_D and R_T explicit is essential: a permission-bearing API operation may belong to R_D without producing privacy data, and is therefore not automatically promoted into R_T .

API identity resolution. The first task is to compute a semantic identity for each candidate API usage. We define:

$$Id(v) = \langle Pkg, Ns, Recv, Mem, Acc \rangle, \quad (4)$$

where Pkg and Ns are resolved through a namespace-constrained migration map, $Recv$ traces the receiver object back to its source via bounded backward data-flow (e.g., `cameraMgr ← getCameraMgr(ctx)`), Mem is the method or property name, and $Acc \in \{direct, assign, indirect, property\}$ classifies the access pattern. Resolution also returns a witness

$$w \in \{Namespace, IRAlias, ExactType, CallTarget, FactoryOrigin, PropertyRead\}. \quad (5)$$

A usage u is recognized when:

$$\exists s \in Stmt(P), \exists r_D \in R_D : Accept(s, r_D, I, A, w), \quad (6)$$

where A is the alias relation recovered from imports, assignments, and IR-lowered namespace-member expressions. The recognized usage set is $U_D = \{u \mid \exists s, r_D, w : Accept(s, r_D, I, A, w)\}$.

Source carriers and seeds. For a statement s matching r_T , the initial taint fact is determined by the carrier rather than by the API name alone:

$$Seed(s, r_T) = \begin{cases} lhs(s), & \text{carrier} = \text{return}, \\ arg_{i_{src}}(s), & \text{carrier} = \text{arg}, \\ param_{i_{src}}(ResolveCb(arg_{i_{cb}}(s))), & \text{carrier} = \text{callback}. \end{cases} \quad (7)$$

Property occurrences in U_D taint the right-hand-side value for detector-local tracing, but enter IFDS only when a corresponding well-formed R_T rule exists. Framework-provided Want values are query sources of kind *framework_input*, not privacy-API detector occurrences. These distinctions prevent detector coverage, privacy-data seeding, and framework-input analysis from being silently equated.

Asynchronous carrier recovery. The second task is to bind a configured asynchronous source to its consumer. We define:

$$SrcBind(call, cb, param), \quad (8)$$

where $call$ is the source API invoke statement, cb is the resolved callback method or Promise handler, and $param$ is the specific parameter through which sensitive data flows. The implementation distinguishes native IFDS propagation from bounded post-IFDS carrier recovery and from callback edges used only for call-chain context; provenance records that distinction.

Provenance-annotated subgraph. The third task is to construct, for each $u \in U_D$, an information-flow subgraph:

$$G_u = \langle V_u, E_u, \lambda_u, \omega_u \rangle, \quad (9)$$

where V_u contains API, caller, sink, and taint-fact nodes; E_u contains call, detector-local sink, configured-query path, and endpoint-link edges. The labeling function λ_u records the semantic edge role. The origin function ω_u is reconstructed from the producing record:

$$\omega_u(e) = \begin{cases} callType(e), & e \in E_{call}, \\ detector_local, & e \in E_{sink}, \\ \tau(q), & e \in E_{path}(q), \\ linkEvidence(e), & e \in E_{link}, \end{cases} \quad (10)$$

where call types distinguish direct, callback, and lifecycle-implicit context; $\tau(q)$ distinguishes IFDS, asynchronous-supplement, and corroborated paths; and link evidence distinguishes exact-statement from normalized-source-location joins. This representation preserves how each subgraph relation was produced without treating derivation type as a probability of correctness.

3.2 Architecture Overview

Figure 2 gives the end-to-end architecture. Project and rule inputs enter ArkAnalyzer for Scene, IR, and call-graph construction. ARKPRISM then performs three connected analyses: (1) API identity resolution computing $Id(v)$ and its witness, (2) carrier propagation across IR and selected asynchronous continuations, and (3) subgraph assembly that preserves the derivation label $\omega(e)$ on




Fig. 2. ARKPRISM analysis pipeline and evidence assembly.

every evidence edge. Lifecycle ordering and call-graph context supply reachability evidence to this pipeline.

3.3 Running Example: Why Direct Calls Are Insufficient

Figure 3 shows an anonymized pattern distilled from manually reviewed benchmark projects. The example is intentionally small, but each highlighted source position corresponds to a real source-level evidence class seen during manual annotation: a manager object returned by a Harmony API, a member call on that manager, a compound namespace API that may be lowered through a temporary, and a local sink that helps an auditor inspect the behavior.

A direct-call baseline that only matches `namespace.method(...)` statements can recover neither `this.calendarMgr.getCalendar(...)` nor `agent.create(...)` reliably. The former is a member call whose receiver obtains its privacy meaning from `calendarManager.getCalendarManager`; the latter is a separate compound-namespace example whose `request.agent` prefix may be assigned to an IR temporary before `create` is invoked. ARKPRISM treats both cases as rule- and evidence-matching problems over imports, aliases, receiver origins, and IR-lowered expressions, then retains lifecycle context, local sink evidence, configured-query paths, and source metadata instead of returning flat method-name hits.

3.4 Namespace-constrained Package Migration

OpenHarmony API evolution makes exact package matching brittle, but consolidated kit packages are not globally equivalent to their historical modules. We therefore define a directional migration relation over package-namespace pairs, $\mathcal{M} \subseteq (Pkg \times Ns) \times (Pkg \times Ns)$. For example, the `deviceInfo` export maps from `@ohos.deviceInfo` to `@kit.BasicServicesKit`, while unrelated exports in the same kit do not become compatible. A detector rule r_D and import binding i are compatible when:

$$Compat(i, r_D) \triangleq (i.pkg, i.ns) = (r_D.pkg, r_D.ns) \vee ((i.pkg, i.ns), (r_D.pkg, r_D.ns)) \in \mathcal{M}. \quad (11)$$




Fig. 3. Evidence-carrying resolution for manager receivers and compound namespaces.

The analyzer also derives namespace aliases from import declarations. If a source file contains `import wifi from '@ohos.wifiManager'`, then $(wifi, wifiManager) \in A_{import}$. If it imports a kit package with a default binding, ARKPRISM records both the textual alias and the canonical namespace. Matching therefore uses a normalized API key:

$$key(u) = \langle \mu(pkg_u, ns_u), \sigma(mem_u) \rangle, \quad (12)$$

where μ applies the migration/export map and import-alias normalization to a package-namespace pair, and σ normalizes method/property spelling. The report retains both $key(u)$ and the raw package signature so that normalized and SDK-specific identities remain auditable.

3.5 Package- and Receiver-aware API Identity Resolution

HapFlow's released namespace-method loader resolves configured entries to SDK declarations, but its schema does not encode executable property reads or retain factory-origin evidence for manager/helper receivers. ARKPRISM represents these access forms by computing a semantic identity and an acceptance witness for each candidate usage.

3.5.1 Identity Function. Given a candidate statement s , ARKPRISM computes $Id(s)$ as defined in Section 3.1. The key components are $Recv(s)$ (receiver-origin tracing, detailed below) and $Acc(s)$ (access-kind classification). The rule mode is normalized as

$$Mode(r_D) = \begin{cases} \{direct, assign\}, & r_D.direct = true, \\ \{indirect\}, & r_D.direct = false, \\ \{property\}, & r_D.direct = null. \end{cases} \quad (13)$$

The candidate predicate is:

$$\begin{aligned}
Match(s, r_D, I, A) &= Compat(I_s, r_D) \\
&\wedge Id(s).Mem = r_D.mem \\
&\wedge Acc(s) \in Mode(r_D).
\end{aligned} \tag{14}$$

Package compatibility only generates candidates. Acceptance additionally requires access-specific evidence:

$$\begin{aligned}
Accept(s, r_D, I, A, w) &\triangleq Match(s, r_D, I, A) \\
&\wedge Sufficient(Acc(s), w, r_D),
\end{aligned} \tag{15}$$

$$\mathcal{W}_{ind} = \{ExactType, CallTarget, FactoryOrigin, Namespace\}. \tag{16}$$

$$Sufficient(a, w, r_D) = \begin{cases} w \in \{Namespace, IRAlias\}, & a \in \{direct, assign\}, \\ w = PropertyRead, & a = property, \\ w \in \mathcal{W}_{ind}, & a = indirect. \end{cases} \tag{17}$$

For indirect calls, *Namespace* means an exact namespace receiver, not merely package presence. It is disabled for generic members such as `get`, `on`, `start`, `request`, and `close`; these require type, target, or factory-origin evidence. Property acceptance requires both a property rule ($r_D.direct = null$) and an executable IR field read. Thus, package import or final-token agreement alone is never sufficient.

IFDS source-rule acceptance. The IFDS query uses SDK method signatures whenever ArkAnalyzer resolves a call target. For a call whose signature remains unknown after frontend recovery, matching falls back to the configured rule metadata rather than to method name alone. Let $Meta(r_T)$ retain the rule's module, namespace, declaring class, member, arity, parameter types, and any string-dispatched event literal. We define:

$$Exact_T(s, r_T) \triangleq TargetSig(s) = Sig(r_T), \tag{18}$$

$$\begin{aligned}
Fuzzy_T(s, r_T) &\triangleq UnkSig(s) \wedge Mem(s) = r_T.mem \\
&\wedge ArityCompat(s, r_T) \wedge ParamCompat(s, r_T).
\end{aligned} \tag{19}$$

The guard for a generic member and the guard for a string-dispatched overload are:

$$\begin{aligned}
IdentityGuard(s, r_T) &\triangleq \neg Generic(r_T.mem) \\
&\vee ReceiverOrModuleWitness(s, Meta(r_T)),
\end{aligned} \tag{20}$$

$$\begin{aligned}
LiteralGuard(s, r_T) &\triangleq \neg HasEventLiteral(r_T) \\
&\vee ArgLiteral(s, 0) = EventLiteral(r_T).
\end{aligned} \tag{21}$$

The source call is accepted iff

$$Accept_T(s, r_T) \triangleq Exact_T(s, r_T) \vee (Fuzzy_T(s, r_T) \wedge IdentityGuard(s, r_T) \wedge LiteralGuard(s, r_T)). \tag{22}$$

Consequently, an unresolved call such as `avPlayer.on('error', ...)` cannot match a configured geolocation or sensor event merely because both use `on` with the same arity. This predicate governs IFDS source seeding and is separate from the detector acceptance predicate above.

Table 1. API identity resolution examples.

Source pattern	$Id(s)$	Acc
<code>id.getOAID(cb)</code>	$\langle \text{BasicSvc}, id, \perp, \text{getOAID}, \text{direct} \rangle$	direct
<code>net=getDefaultNet()</code>	$\langle \text{NetworkKit}, net, \perp, \text{getDefaultNet}, \text{assign} \rangle$	assign
<code>mgr.getData(cb)</code>	$\langle \text{BasicSvc}, mgr, \text{getMgr}, \text{getData}, \text{indirect} \rangle$	indirect
<code>deviceInfo.brand</code>	$\langle \text{BasicSvc}, devInfo, \perp, \text{brand}, \text{property} \rangle$	property

3.5.2 Receiver-Origin Tracing. The key addition beyond the released namespace–method loader is $Recv(s)$. For a statement $s : x.m(\dots)$ where x is a local or field variable, ARKPRISM traces x backward through assignments and instance-invoke bases to find its origin, subject to a tracing bound of k_{recv} steps (default 8) to avoid unbounded backward exploration:

$$Recv(s) = \begin{cases} Ns(def), & x \leftarrow \text{invoke}(q.f()), f \in \text{factory}(r_D), \\ Recv(def), & x \leftarrow y, y \text{ has known origin,} \\ TF(x), & \text{type}(x) = T, T \in SDK, \\ \perp, & \text{unresolvable in } k_{recv} \text{ steps.} \end{cases} \quad (23)$$

The first case covers the common pattern $mgr \leftarrow \text{getManager}(ctx); mgr.getData(cb)$, where the receiver mgr originates from a privacy API factory call. The second case propagates origin through intra-procedural assignment chains and instance-invoke bases. The third case (TF) uses an exactly compatible declared SDK type. Inter-method field assignments are not claimed as origin-proven unless the defining statement or an exact type/target identity is recovered; otherwise the call is rejected. This rule avoids silently treating package import plus method-name agreement as a typed call target.

3.5.3 Access-Kind Classification. Each recognized usage receives an Acc label:

- *direct*: bare invoke statement, e.g., `id.getOAID(cb)`;
- *assign*: invoke with return value assigned, e.g., `net = getDefaultNet()`;
- *indirect*: invoke on a traced receiver, e.g., `mgr.getData(cb)` where $Recv(mgr) \neq \perp$;
- *property*: field/property read, e.g., `deviceInfo.productModel`.

Table 1 illustrates how Id resolves each access kind.

IR-lowered multi-segment APIs require an additional alias relation A_{ir} . For assignments of the form $t := q.member$, ARKPRISM records $(t, q.member) \in A_{ir}$. A later statement $t.create(\dots)$ is matched as $q.member.create(\dots)$ if the composed signature appears in R_D . This is essential for APIs such as `request.agent.create`, whose namespace can disappear after lowering.

3.6 Lifecycle State Machine Model

OpenHarmony applications follow a structured lifecycle managed by the system runtime. A UIAbility transitions through creation, foreground, background, and destruction phases; custom components follow appearance, build, page visibility, and disappearance phases; UI callbacks are registered during component build and triggered by user interaction. These transitions are *implicit*—the runtime invokes lifecycle methods without visible call edges in application code—so a static analyzer that treats each lifecycle method as an independent entry point misses cross-phase data flow.

Table 2. Asynchronous carrier and context recovery rules.

Rule	Construct	Recovery relation	Output	Domain
T1	Callback source	$Resolve(a_i) \rightarrow cb.param_j$	May-path	Post-IFDS
T2	Field callback	$Read(this.f) \rightarrow cb$	Call edge	Report CG
T3	Builder option	$Arg(bld, f) \rightarrow cb$	Call edge	Report CG
T4	Promise then	$ret(src) \rightarrow then_j.param$	May-path	Post-IFDS
T5	async/await	$resolve_arg \rightarrow await_{res}$	May-path	Post-IFDS

ARKPRISM models the OpenHarmony application lifecycle as a three-layer state machine. Let the lifecycle model be:

$$\mathcal{L} = \langle \Phi_a, \Phi_c, \delta_a, \delta_c, \delta_{cross} \rangle, \quad (24)$$

where Φ_a is the set of ability phases (e.g., CREATING, FOREGROUNDED), Φ_c the component phases (e.g., APPEARING, BUILT), $\delta_a \subseteq \Phi_a \times \Phi_a$ the ability transition relation, $\delta_c \subseteq \Phi_c \times \Phi_c$ the component transition relation, and $\delta_{cross} \subseteq \Phi_a \times \Phi_c$ the cross-layer transition relation. Cross-layer transitions capture implicit calls such as `onForeground` $\rightarrow$ `aboutToAppear`.

The lifecycle modeler discovers ability classes by checking superclass inheritance from `UIAbility`, `Ability`, and extension ability base classes; component classes by checking superclass (`CustomComponent`, `ViewPU`), decorators (`@Component`), or the presence of two or more component lifecycle methods; and callback methods by scanning IR statements for `CommonMethod.*` invocations and extracting `FunctionType` arguments.

Lifecycle information is consumed in two distinct ways. For call-chain recovery, each transition $(m_i, m_j) \in \delta_a \cup \delta_c \cup \delta_{cross}$ becomes a provenance-labeled implicit edge. For IFDS, the model orders entry methods in a synthetic `DummyMain`: ability creation and foreground callbacks precede component initialization and UI callbacks, followed by background and destruction phases. The bounded default removes the standard back edge from the final entry to the loop header, preventing arbitrary destruction-to-creation cycles. Repeatable event classes may receive one additional explicit firing, but no unbounded loop. The IFDS solver derives interprocedural targets from the Scene and pointer information; it does not consume the separately augmented report call graph.

3.7 Asynchronous Carrier and Context Recovery

OpenHarmony privacy APIs often deliver sensitive data through asynchronous callbacks and Promise chains rather than direct return values. For example, `geoLocationManager.getCurrentLocation(callback)` passes location data to a callback parameter, while `wifiManager.getLinkedInfo().then(callback)` delivers Wi-Fi state through a Promise resolution. Standard call-to-return propagation does not by itself establish the connection between a source invocation and a callback parameter or Promise resolution, so these constructs require explicit transfer rules.

ARKPRISM formalizes carrier binding as $SrcBind(call, cb, param)$ (Section 3.1). The five recovery rules do not all modify the same analysis. T1, T4, and T5 add bounded source-to-sink outcomes after the core IFDS run when ArkIR lacks an explicit continuation edge; T2 and T3 add callback edges to the report call graph and therefore improve context recovery rather than seeding IFDS facts. Table 2 makes this execution domain explicit.

3.7.1 T1: Callback Parameter Binding. When an R_T source accepts a callback as a direct argument at position i , the binding is $SrcBind(s, cb, param_j)$ where $cb = Resolve(a_i)$ and $param_j$ is the callback

parameter selected by the rule's carrier metadata. The post-IFDS phase accepts only typed or definition-backed callbacks. Resolution applies three strategies in order:

$$\text{Resolve}(a_i) = \begin{cases} m_f, & a_i \in \text{FuncType} \wedge m_f = \text{Lkp}(a_i.\text{sig}) \\ m_c, & a_i \in \text{ClosureType} \wedge m_c = \text{Lkp}(\text{Ext}(a_i)) \\ m_l, & a_i \in \text{Local} \wedge m_l = \text{Trc}(a_i) \end{cases} \quad (25)$$

FunctionType resolution (*Lkp*) extracts the method signature from the compiler-generated function type and looks up the corresponding method in the declaring class. *ClosureType* resolution (*Ext+Lkp*) parses the closure variable name from the closure type string (e.g., closures: __lambda1) and performs a method name lookup. *Local* variable tracing (*Trc*) follows the definition chain: if a_i is assigned from a new `LambdaClass()` expression, the *ClosureType* of the constructor call provides the callback method. Example: `getData((err,data)⇒...)` resolves the arrow function to its IR method and binds `data` as the sensitive parameter.

The tracer seeds only param_j selected by $r_T.\text{isrc}$, not every callback parameter. It propagates through ArkIR value-use edges and assignments, and accepts a sink only when a sink argument structurally depends on the current fact. IR-local names are never compared by substring, and control dependence alone does not constitute an explicit data-flow edge.

3.7.2 T2: Field-Stored Callback Binding. ArkUI components frequently store callback functions in class fields (e.g., `this.handler = () => {...}`). The callback is registered with a framework API at a later point, creating a temporal gap between definition and invocation. ARKPRISM reads the field value at the registration site and traces it to the assignment, adding $\text{Read}(\text{this}.f) \rightarrow cb$ to the report call graph. This handles the pattern where `this.onSelect` is passed to `List.onSelect(this.onSelect)` in a component's build method. T2 supplies entry/call context; it does not by itself assert a sensitive carrier or inject an IFDS fact.

3.7.3 T3: Builder-Option Callback Binding. ArkUI builder APIs accept configuration objects whose fields are callback functions (e.g., `bindPopup({builder: popupBuilder})`). T3 extracts the callback from a named option field, $\text{Arg}(\text{bld}, f) \rightarrow cb$, and adds the recovered owner-builder edge to the report call graph. The pattern is structurally distinct from T1 because the callback is nested inside a lowered object initializer rather than passed as a positional source callback. As with T2, the edge provides call context and is not an IFDS seed.

3.7.4 T4: Promise-then Chain Binding. Promise `.then()` chains can carry a source return value through Promise resolution to a callback parameter. When the core IFDS result lacks that continuation, the bounded supplement recognizes three sub-patterns:

Direct chaining: $\text{sourceApi}(\text{args}).\text{then}(cb)$ creates $\text{SrcBind}(\text{sourceStmt}, cb, \text{param})$ with $\text{Edge}(\text{sourceStmt}, cb_{\text{param}})$.

Variable chaining: $x = \text{sourceApi}(\text{args}); x.\text{then}(cb)$ creates $\text{Edge}(\text{sourceStmt}, x_{\text{assign}}) \cup \text{Edge}(x_{\text{assign}}, cb_{\text{param}})$.

Sequential chaining: $x.\text{then}(cb_1).\text{then}(cb_2)$ creates $\bigcup_i \text{Edge}(cb_{i-1}, cb_i)$, propagating taint through the chain.

3.7.5 T5: Async/Await Bridge Binding. An asynchronous function that awaits a Promise-wrapped source API creates an implicit binding between the resolved value and the local variable:

$$\text{data} = \text{await promiseMethod()} \Rightarrow \text{Edge}(\text{resolve}_{\text{arg}}, \text{await}_{\text{res}}) \cup \text{Edge}(\text{await}_{\text{res}}, \text{data}_{\text{uses}}) \quad (26)$$

where $resolve_{arg}$ traces from the $resolve()$ call inside the Promise executor to the awaited result. The supplement requires a definition-backed Promise owner and separately labels the resulting path; it does not claim that the core IFDS supergraph contains the implicit executor-to-await edge.

3.7.6 Algorithm. Algorithm 1 summarizes the post-IFDS portion (T1, T4, and T5). It does not modify IFDS flow functions or the exploded supergraph. Instead, it scans configured sources whose continuation edge is absent, resolves only typed callbacks or definition-backed Promise owners, performs bounded local/interprocedural tracing, and appends separately provenance-labeled outcomes. T2 and T3 are constructed by the callback-augmented report call graph in Section 3.7; they are intentionally absent from Algorithm 1.

Algorithm 1 Post-IFDS Asynchronous Carrier Recovery

Require: Program Scene P , IFDS source rules R_T , existing outcomes O , bounds B

Ensure: Augmented outcomes O' ; each new outcome preserves $\epsilon(src)$ and has provenance ASYNCSUPPLEMENT

```

1:  $A \leftarrow \emptyset$ 
2: for all application methods  $m \in P$  do
3:   for all invoke statements  $s \in m$  do
4:      $src \leftarrow CALLSOURCE(s, R_T)$ 
5:     if  $src \neq \emptyset \wedge src.carrier = callback$  then ▷ T1
6:        $cb \leftarrow RESOLVE(s.args[src.cbIdx])$ 
7:       if  $cb \neq \emptyset \wedge src.srcIdx < |cb.params|$  then
8:          $p \leftarrow cb.params[src.srcIdx]$ 
9:          $A \leftarrow A \cup TRACEToSINK(cb, p, Path(s), \epsilon(src), B)$ 
10:      end if
11:    end if
12:    if  $s$  is a then invocation on base  $b$  then ▷ T4
13:      if  $b$  is a configured source invocation then
14:         $si \leftarrow b$ 
15:      else
16:         $si \leftarrow FINDSRCDEF(m, b)$ 
17:      end if
18:      if  $si \neq \emptyset$  then
19:         $cb \leftarrow RESOLVE(s.args[0])$ 
20:        if  $cb \neq \emptyset$  then
21:           $A \leftarrow A \cup TRACEThenToSINK(m, si, cb, \epsilon(si), B)$ 
22:        end if
23:      end if
24:    end if
25:  end for
26:  for all await sites  $a \in m$  with promise  $p$  do ▷ T5
27:     $ow \leftarrow FINDPROMISEOWNER(p, R_T)$ 
28:    if  $ow \neq \emptyset$  then
29:       $rf \leftarrow COLLECTRESOLVEFACTS(ow)$ 
30:      if  $rf \neq \emptyset$  then
31:         $A \leftarrow A \cup TRACEToSINK(m, a.res, Path(rf) \cup Path(a), \epsilon(ow), B)$ 
32:      end if
33:    end if
34:  end for
35: end for
36:  $O' \leftarrow DEDUPLICATE(O \cup A)$ 
37: return  $O'$ 

```

3.8 Callback-augmented Call-chain Recovery

Given a recognized usage u inside method m_u , call-chain recovery searches a reverse graph:

$$E_r = (E_c)^{-1} \cup (E_{invoke})^{-1} \cup (E_{cb})^{-1} \cup (E_{init})^{-1}. \quad (27)$$

Here E_c is the ArkAnalyzer call graph, E_{invoke} contains IR-level call edges absent from the graph, E_{cb} contains callback edges recovered from `FunctionType/ClosureType` values, field callback assignments, and builder-option callbacks, and E_{init} links generated initialization methods to their owning class or component context. The recovered chain set is:

$$C_u = \{ \pi = (m_0, \dots, m_u) \mid m_0 \in \text{Entry} \cup \text{Init} \cup \text{Unk}, \\ (m_i, m_{i+1}) \in E_r^{-1} \}. \quad (28)$$

Each chain receives an entry-evidence class:

$$\kappa(\pi) = \begin{cases} \text{framework}, & m_0 \in \text{Entry} \\ \text{initialization}, & m_0 \in \text{Init} \\ \text{local}, & \pi = (m_u) \text{ and no predecessor is found} \\ \text{unk}, & \text{otherwise.} \end{cases} \quad (29)$$

ARKPRISM does not collapse these classes into a single coverage number. A local fallback chain is useful evidence that the API usage exists in a method, but it is not a recovered framework-entry path. This distinction is preserved in the report and in the evaluation tables.

3.9 Sink and Taint Attachment

For each detector occurrence $u \in U_D$, local tracing extracts statements $K_D(u)$ that may expose privacy-related data through logging, UI display, storage, network transfer, or data return. Sink resolution requires an exact configured method token and compatible receiver/namespace evidence; in particular, substring agreement such as `set` versus `setData` is rejected.

The IFDS query is configured independently by R_T and its sink set K_T . Let S_P and S_F be source statements whose rules have kind *privacy_data* and *framework_input*, respectively. The solver computes

$$T_I \subseteq S_T \times K_T \times \text{Path}, \quad (30)$$

where $S_T = S_P \uplus S_F$ is the disjoint typed-source universe resolved from R_T . Each tuple $\langle s, k, \rho \rangle$ means that the carrier seeded at s may reach configured sink k along path ρ . Each seeded fact carries an immutable source witness

$$\epsilon(s) = \langle \text{kind}, \text{pkg}, \text{ns}, \text{cls}, \text{mem}, \text{carrier}, \text{sig}, \text{rule}, \text{loc} \rangle, \quad (31)$$

where $\text{kind} \in \{\text{privacy_data}, \text{framework_input}\}$ and the remaining fields identify the resolved API, carrier position, SDK signature, rule provenance, and normalized source location. Normal, call, return, and exceptional transfers copy $\epsilon(s)$ unchanged; fact and edge identities include it. Consequently, values reaching the same statement from different configured sources remain distinct, and the report obtains its source endpoint from $\epsilon(s)$ rather than inferring it from the first statement in ρ . The bounded asynchronous rules preserve the same witness and produce a separate set

$$T_A \subseteq S_T \times K_T \times \text{Path}, \quad (32)$$

and the report exposes $T_Q = T_I \cup T_A$, stratified as

$$T_P = \{q \in T_Q \mid \text{src}(q) \in S_P\}, \quad T_F = \{q \in T_Q \mid \text{src}(q) \in S_F\}. \quad (33)$$

For duplicate endpoint/path outcomes, path provenance is retained as

$$\tau(q) \in \{ifds, async_supplement, both\}. \quad (34)$$

A partial linkage relation

$$L \subseteq U_D \times T_P \quad (35)$$

is created only when project, normalized source location, and compatible API identity agree. Framework-input paths are never linked to privacy-API detector occurrences. Linked privacy-data paths are attached to the corresponding G_u ; unlinked configured-query paths remain project-level taint evidence rather than being forced onto a detector occurrence. Consequently, detector-local sink-positive projects and configured-query-path-positive projects are not set-inclusion counts. The corpus evaluation reports both evidence universes and both source kinds separately.

3.9.1 Provenance Annotation. The report stores provenance at the level where it is produced. A call-chain relation carries a `callType`; a configured-query path carries $\tau(q) \in \{ifds, async_supplement, both\}$; and an endpoint link records `exact_statement` or `source_location`. Detector-local sink edges are identified by their containing record. Subgraph assembly maps these fields to ω_u rather than replacing them with a synthetic global confidence score. Consequently, an asynchronous supplement cannot be presented as an IFDS path, and a normalized-location join remains distinguishable from an exact-statement join.

3.10 Analysis Contract

ARKPRISM's recognition claim is parameterized by three explicit inputs: ArkAnalyzer IR for the source construct, the fixed API rule set, and the namespace, receiver, origin, or target-signature evidence required by *Accept*. The analysis targets statically represented ArkTS code; native code, dynamic loading, reflection-like runtime dispatch, and APIs absent from the rule set are outside the query domain. A recognized API establishes an executable use under this contract, while policy-level violation assessment remains a downstream audit decision.

IFDS and asynchronous-supplement results are may-flows under their respective source/sink and reachability models. The report preserves their producer and does not convert a detector occurrence into an IFDS source without a source-endpoint link. Every run records one of three solver states: complete, resource-bounded, or failed. Resource-bounded and failed cases are never reclassified as zero-flow applications. Edge deduplication is scoped by IR statement, value identity, access path, and source witness; textual equality alone is insufficient because same-named locals in different methods and facts from different seeds are distinct. These requirements are regression-tested on HapBench positive and negative cases and are also enforced by the isolated-process corpus runner.

3.11 Scalable Execution

SDK files are loaded lazily into the Scene when taint analysis starts, keeping frontend construction independent of the query-specific SDK augmentation. Source and sink rules expand SDK-path placeholders at runtime, and IFDS resource bounds are explicit. Each project runs in a fresh process, while per-project logs and status files preserve timeout and failure states. If pointer analysis fails, the runner retains only evidence from completed producing phases. This execution model prevents incomplete states from entering the aggregate as negative results.

Let C be the number of executable call/property candidates, d the maximum number of package/namespace-compatible rules returned by the rule index, and k_{recv} the receiver-definition bound. Identity resolution costs $O(Cdk_{recv})$ time and $O(C)$ output space; in practice d is small because package and namespace indexing precede member matching. The evidence carried by a source fact is immutable and constant-sized, so it refines fact identity without changing the

underlying IFDS formulation. The post-IFDS phase is independently capped by the number of source candidates, visited methods and values, emitted outcomes, and path length. It therefore cannot expand into an unbounded event-loop search; paths beyond these bounds are not emitted, and every retained outcome remains labeled as supplementary rather than as an IFDS proof. These bounds make the extra analyses predictable while preserving the provenance needed to distinguish the two producers.

The complete pipeline follows a fixed order: Scene construction, lifecycle modeling, report call-graph construction, import/alias extraction, API identity resolution via *Accept*, callback-augmented chain recovery, detector-local sink tracing, IFDS taint analysis from the lifecycle-aware *DummyMain*, bounded asynchronous supplementation, endpoint linkage, and subgraph assembly. A key property is that ARKPRISM never treats a missing layer as proof of absence: if a call-chain predecessor is absent, the usage still becomes a local evidence chain; if a taint path has no compatible detector endpoint, both records remain visible but unlinked.

4 Implementation

ARKPRISM is implemented in TypeScript and runs as a command-line analyzer. It reuses ArkAnalyzer's Scene, IR, and call-graph abstractions, and integrates HapFlow through a runner module.

Rule bases. The release configuration in `sensitive_apis.json` contains 58 nonempty system-package groups and 716 detector records (710 unique package-namespace-member identities and 665 unique namespace-member identities): 453 direct, 213 manager/receiver, and 50 property records. Each record stores namespace, method/property, permission, privacy category, access mode, and optional receiver factories. The taxonomy contains 38 privacy categories and 74 distinct permission labels. Deterministic integrity tests enforce required fields and access modes, while output canonicalization and stable rule hashes keep repeated configuration records from inflating reported occurrences.

Rule construction. The detector inventory is a frozen policy input rather than a set mined from ARKPRISM output. Construction retains its package, namespace, member/property, permission, category, and access-mode decisions, then checks structural validity and namespace-constrained SDK migration. The taint-source universe applies a stricter carrier protocol. Starting from 913 candidate source rules, an auditable migration rejects 655 entries without a supported source kind, concrete non-void return, typed callback parameter, or valid carrier index. Exact API-20 declaration resolution and parameter normalization then produce the 257-rule privacy-data inventory used in the experiments. The migration artifact retains every rejected rule and reason; SDK audits retain unresolved, ambiguous, version-incompatible, and carrier-mismatch records. This process makes the evaluated rule contract reproducible without treating the resulting taxonomy as an exhaustive definition of privacy sensitivity.

Detector coverage and taint-source seeding use separate rule universes. The IFDS inventory admits 257 privacy-data rules whose API-20 declarations expose a concrete return or callback carrier, plus two explicitly modeled *framework_input* lifecycle rules. Setters, publication/network actions, void/any returns, external SDK owners, and executable properties without a modeled field carrier remain detector evidence rather than becoming speculative taint seeds. This carrier contract keeps permission-bearing operations visible for auditing while ensuring that every propagated fact names the value that may contain sensitive data.

The resolver treats an explicit parameter array, including the empty array, as an exact arity constraint and compares normalized parameter types and literal discriminators. This prevents a zero-argument Promise rule from absorbing a callback overload and prevents event-specific on/off rules from matching another event with the same parameter names. Module, namespace, and class are hard owner constraints: if an explicitly named class is unavailable in the loaded SDK, resolution

fails instead of falling back to another class with the same member name. Against the fixed API-20 SDK, all 257 privacy-data rules resolve to exactly one declaration: there are no unresolved or ambiguous rules, signature collisions, or carrier-type mismatches. The two framework-input rules likewise resolve uniquely under both the API-17 benchmark SDK and API-20 corpus SDK. The API-17 compatibility audit resolves 249 of the 257 privacy-data rules and records the eight version-incompatible rules without substituting a different owner. The loader also rejects a missing source kind, non-concrete return carrier, callback type, or carrier index. The structural audit reports no malformed detector, source, or sink entries and 29 configured sinks. A namespace-constrained migration table maps historical modules to exports of consolidated `@kit.*` packages; it is not treated as whole-package equivalence. Detector rules, privacy-data sources, framework inputs, SDK-resolution records, and hashes are retained separately.

Detector. The API detector operates per Ark file: it obtains system imports, expands package aliases, divides rules into direct, indirect, and constant-like sets, and scans every method CFG. Method matching normalizes `Class.method`, `method()`, and event-style signatures while preserving namespace context. Package-compatible rules form the candidate set; indirect calls are then accepted only with an exact SDK receiver type, exact declaring-class identity from the call-target signature, an intra-procedural definition/factory-origin chain (bounded at eight definitions), or an exact namespace receiver for a non-generic member. The selected evidence is stored in `matchEvidence`; package-only and ambiguous method-only matches are rejected. For IR-lowered multi-segment APIs (e.g., `request.agent.create`), the detector records namespace-member aliases when ArkAnalyzer lowers the namespace prefix into a temporary.

Call-chain tracer. The tracer builds an enhanced reverse call map from ArkAnalyzer call-graph edges plus supplementary edges recovered from `FunctionType/ClosureType` signatures, class-field callback assignments, and ArkUI builder-option callbacks. Entry methods are classified through lifecycle and callback naming conventions; if no upstream entry is found, the tracer emits a local fallback chain.

Detector-local sink analysis. For each detector occurrence, the sink analyzer follows local and callback-related values and classifies observations into log, UI display, storage, network, data return, intent, and share categories. A semantic-enrichment pass adds guard conditions, try-catch context, and loop evidence. These observations are not the IFDS sink universe.

IFDS runner and asynchronous supplement. The runner lazily loads SDK files, creates the lifecycle-aware `DummyMain` entry (Section 3.6), attempts pointer analysis, and loads the independent IFDS source/sink query. The solver derives call targets from Scene and pointer-analysis evidence, including receiver-refined indirect targets; it does not reuse the separately augmented report call graph. Unknown-signature source matching retains the rule's module, namespace, class, arity, and parameter metadata. Generic members such as `on`, `get`, and `create` require receiver/module evidence; a string-dispatched overload such as `on('locationChange', ...)` must additionally match the configured event literal. Thus, an unrelated `avPlayer.on('error', ...)` cannot inherit a geolocation or sensor source solely from the member name and arity.

IFDS edge keys combine IR object identity, access-path state, and immutable source evidence rather than statement text, preventing distinct same-named locals or different source seeds from being merged. Normal, call, return, exceptional, callback, Promise, and `await` propagation preserve this evidence. Exception transfer is value-dependent: an Error payload is tainted only when one of its constructor arguments depends on the incoming fact. The solver runs all sources together by default; batching is available only as an explicit resource fallback. After IFDS terminates, a separately bounded callback/Promise scan recovers typed callback, then, and `await` paths that are absent from the supergraph. It seeds the configured callback-parameter index and uses structured `Value.getUses/ValueEqual` dependencies rather than textual IR-local matching; in particular,

temporary %8 does not match %85. Outcomes are deduplicated only when their source evidence, tainted values, and complete statement paths agree, then merged with, but not relabeled as, solver outcomes. Unknown-signature sinks likewise require method and receiver/namespace agreement, and failure or budget exhaustion is recorded rather than emitted as an empty negative result.

Evidence linkage. Every path records whether its seed is *privacy_data* or *framework_input*, together with its resolved module, namespace, class, member, carrier, SDK signature, and rule origin. Subgraph assembly considers only privacy-data paths, and joins one to a detector occurrence only when project, normalized source location, and compatible API identity agree. The release runner rejects a report whose path lacks source evidence or whose first normalized endpoint disagrees with that evidence. Detector-local sink observations, unlinked privacy-data paths, and framework-input paths remain separate records. This explicit partial join prevents aggregate IFDS paths from being interpreted as downstream sinks of every detector occurrence.

Lifecycle and call-graph context. The lifecycle modeler builds $\mathcal{L}$ (Section 3.6) and orders ability, component, and event entries in *DummyMain*. The default removes unbounded lifecycle back edges while allowing an explicit second firing for selected repeatable events. Separately, the report call graph augments RTA/CHA with provenance-labeled lifecycle transitions, while *ArkAnalyzer*'s pointer analysis supplies receiver-sensitive targets to IFDS.

5 Evaluation

We answer four questions: **RQ1** How accurately does *ARKPRISM* resolve privacy-API identities, and which evidence forms account for that accuracy? **RQ2** How accurately do *ARKPRISM* and the author-released *HapFlow* artifact classify source-to-sink flows? **RQ3** How do IR recovery, receiver refinement, lifecycle bounds, and asynchronous recovery affect analysis quality? **RQ4** What evidence and scalability patterns appear across a large *ArkTS* corpus? Each oracle supports a distinct claim: source-first benchmarks evaluate identity, *HapBench* evaluates executable flow classification, *ArkAsyncBench* attributes asynchronous coverage, a stratified source audit evaluates real-project path semantics, and the corpus measures prevalence and cost.

5.1 Experimental Design

Implementations. We evaluate the released *ARKPRISM* implementation and the author-released *HapFlow* artifact. Both use *ArkAnalyzer* IR and the same source-sink query on *HapBench*. That reproduction uses the API-17 SDK bundled with *HapFlow*; the real-project benchmarks and corpus use API-20 at `E:\OpenHarmony_SDK\20\ets`. SDK fallback is disabled. The corpus runner uses one analyzer process per project and two workers, preventing heap state or failure in one project from contaminating another. IFDS runs once per project without batching.

Identity benchmarks. *ArkSourceFirst60* fixes ten projects from each of six project-scale/configured-import strata before detector execution. A project-wide AST pass enumerates candidates but does not label them; source review then assigns 186 exact-location decisions. The resulting oracle contains 90 sensitive occurrences in 16 projects, 96 same-name or owner-incompatible negatives, and 44 confirmed-zero projects. *ArkIdentityBench* contains 32 positive and 32 matched-negative cases spanning imports, typed and constructed receivers, factories, properties, compound namespaces, and package migration.

The *Top-120 stress audit* is deliberately output-selected. Its 1,533 reviewed location rows canonicalize to 666 confirmed project-API keys with executable source evidence. The final detector recovers all 666 keys and reports 182 additional keys outside that frozen audit; we therefore report gold-key reproduction coverage, not precision over the expanded output.

Flow benchmarks. *HapBench* comprises all 67 executable cases released with *HapFlow*: 53 positive and 14 negative. We preserve its published case oracle and additionally audit every canonical

source-sink endpoint pair reported by ARKPRISM. *ArkAsyncBench* contains 48 source-first cases: 24 positives and 24 matched negatives over callback (T1), Promise-then (T4), and `await` (T5) carriers. Positives pass the configured carrier, a carrier property, or an alias to a sink; negatives pass an error, constant, rejection value, ignored result, or independent same-named local.

Real-project path audit. From the final 449-path corpus, a deterministic diversity sampler fixes 120 paths from 48 projects: 90 privacy-data and 30 framework-input paths; 64 IFDS, 30 supplementary, and 26 dual-provenance paths; and 23 short, 70 medium, and 27 long paths. The sample includes every network and UI path stratum. Each record is bound to the run manifest and source/report hashes. Source and IR review independently judges source identity, sink identity, explicit value dependence, reachability, and provenance; a path is correct only when all five dimensions hold.

Large corpus. The corpus contains 1,014 ArkTS projects analyzed under one frozen build, SDK, rule set, and resource contract. It covers applications, libraries, and capability demonstrations; the experiment therefore characterizes this collection rather than a marketplace. The manifest records all input hashes and one verified report per project. Corpus aggregation rejects report-count mismatches, invalid endpoints or provenance, unresolved trace locations, and resource-bounded results.

Metrics. Identity occurrences match on project, normalized file, line/column, canonical package/-namespace/member, and access kind. Flow benchmarks report precision, recall, specificity, balanced accuracy, F1, and MCC. We use Wilson intervals for binomial rates and exact two-sided McNemar tests for paired predictions. Corpus statistics report project prevalence, concentration (Gini and HHI), Spearman associations, source scope, clone sensitivity, and runtime; they are descriptive rather than accuracy estimates.

5.2 RQ1: Privacy-API Identity

Source-first accuracy. ARKPRISM reports exactly the 90 gold occurrences in ArkSourceFirst60: TP/FP/FN= 90/0/0, for 100% observed precision, recall, and F1 (95% Wilson interval: 95.91–100%). It also classifies all 16 positive and 44 zero projects correctly. Exact-location matching prevents one call from masking a missed occurrence elsewhere in the same file.

The gold set contains 44 executable namespace properties, 26 namespace calls, 11 typed receivers, four constructed receivers, three manager receivers, and two factory receivers. Thus, 20/90 occurrences require receiver evidence beyond the accepted line's namespace syntax, and nearly half are property reads rather than calls. The 96 negatives stress the complementary boundary: application wrappers, UI controllers, collections, test drivers, and third-party methods reuse generic members such as `create`, `on`, `request`, and `getData`.

Evidence progression. Table 3 shows how the decision boundary changes. Lexical matching reaches every positive but accepts every matched collision. Import compatibility removes these collisions; receiver and factory evidence then recover indirect uses. The complete resolver adds constructed receivers, executable properties, compound aliases, and namespace-constrained migration, classifying all 64 cases correctly.

High-volume reproduction. The final configuration recovers all 666 manually confirmed project-API keys in the Top-120 stress audit (100% reproduction coverage; 95% Wilson lower bound 99.43%), including 558 namespace-qualified, 102 member/receiver, five template-expression, and one compound-qualified key. The 182 new keys remain outside this frozen oracle and are not folded into a precision numerator or denominator.

Answer to RQ1. On an output-independent real-project oracle, ARKPRISM resolves all 90 sensitive occurrences and rejects all 96 candidate collisions. The controlled benchmark attributes this result

Table 3. Identity evidence progression.

Configuration	TP	TN	FP	FN	Prec.	Rec.	Spec.	F1
Lexical token	32	0	32	0	50.00	100.00	0.00	66.67
Import-aware namespace	17	32	0	15	100.00	53.13	100.00	69.39
Declared receiver	26	32	0	6	100.00	81.25	100.00	89.66
Factory-aware receiver	28	32	0	4	100.00	87.50	100.00	93.33
ARKPRISM	32	32	0	0	100.00	100.00	100.00	100.00

Table 4. HapBench case-level classification.

Tool	TP	TN	FP	FN	Precision	Recall	Specificity	F1	BAcc.	MCC
HapFlow [15]	51	10	4	2	92.73%	96.23%	71.43%	94.44%	83.83%	0.717
ARKPRISM	50	14	0	3	100.00%	94.34%	100.00%	97.09%	97.17%	0.881

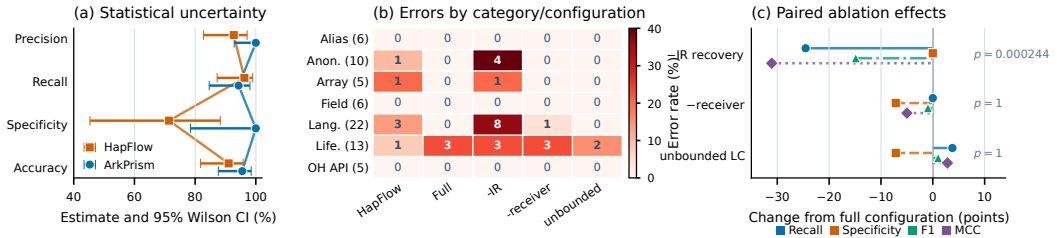


Fig. 4. HapBench uncertainty, category errors, and paired ablation effects.

to the composition of package, receiver, factory, property, and IR evidence; the stress audit shows that the final configuration preserves all 666 previously confirmed high-volume keys.

5.3 RQ2: End-to-End Flow Classification

Table 4 reports the complete HapBench reproduction. ARKPRISM classifies 64/67 cases correctly (50 TP, 14 TN), while HapFlow classifies 61/67 (51 TP, 10 TN). ARKPRISM rejects every negative case and retains 50/53 positives, yielding 100% precision and specificity, 97.09% F1, 97.17% balanced accuracy, and 0.881 MCC.

The tools disagree on nine cases: ARKPRISM alone is correct on six and HapFlow alone on three (exact paired $p = 0.508$). The observed advantage is concentrated in negative-case discrimination. Figure 4(a) shows Wilson uncertainty, and panel (b) localizes every error across seven categories and five configurations. ARKPRISM reaches 100% F1 in six categories and 85.71% in lifecycle cases.

The endpoint-pair audit tests precision inside positive cases. The 52 raw paths canonicalize to 51 source-sink pairs: 50 match the benchmark behavior, while one additional virtual target appears inside a positive case. With three missing positive endpoints, pair-level TP/FP/FN is 50/1/3, giving 98.04% precision, 94.34% recall, and 96.15% F1.

Where the classifications differ. ArkTS-specific IR recovery resolves the anonymous initializer and exceptional successor missed by HapFlow, while constant-index access paths, receiver-refined dispatch, and bounded lifecycle transitions remove its array, virtual-target, and cross-event false positives. The three published positives not emitted by ARKPRISM require, respectively, an implicit

Table 5. HapBench ablation results.

Configuration	TP	TN	FP	FN	Recall	Specificity	BAcc.
ARKPRISM	50	14	0	3	94.34	100.00	97.17
–IR recovery	37	14	0	16	69.81	100.00	84.91
–receiver refinement	50	13	1	3	94.34	92.86	93.60
Unbounded lifecycle	52	13	1	1	98.11	92.86	95.49

Table 6. Asynchronous carrier provenance on ArkAsyncBench.

Construct	Pos./Neg.	IFDS	Supplement	Dual	Supp.-only	Real audit
T1 callback	9/9	9		9	0	1/1
T4 then	9/9	0		9	9	1/1
T5 await	6/6	6		6	0	6/6
Total	24/24	15		24	9	8/8

parent lifecycle execution, state continuity across framework-created extension instances, or a later ability creation after a click callback. Appendix A.2 evaluates this event-order distinction without changing the primary published oracle.

Answer to RQ2. ARKPRISM combines 100% case-level precision and specificity with 97.09% F1 on the full executable suite. Its principal empirical advantage over the released HapFlow artifact is precise rejection of negative cases; the endpoint audit retains 98.04% precision at the finer source–sink-pair unit.

5.4 RQ3: Mechanism and Path Quality

IR, receiver, and lifecycle mechanisms. Table 5 reports pre-specified ablations under the same query. Removing IR recovery loses 13 additional positives; the full configuration is exclusively correct on 13 cases and never exclusively wrong ($p = 0.000244$, significant after Bonferroni correction for four tests). Receiver refinement removes one infeasible virtual target without reducing recall.

The bounded default preserves all 14 negatives and attains 97.17% balanced accuracy, compared with 95.49% under unbounded lifecycle cycling. Removing the bound accepts two additional published positives but also recreates the benchmark’s explicit unreachable-flow negative. Because F1 omits true negatives, it ranks the unbounded setting higher on that narrower objective; balanced accuracy measures the positive and negative obligations together. Under the strict executed-order contract in Appendix A.2, bounded ARKPRISM classifies 67/67 cases, whereas unbounded cycling creates three cross-event false positives. The default therefore targets the audit objective: preserve negative-case discrimination unless an executable transition witness exists.

Asynchronous provenance. HapBench exercises callbacks with explicit IFDS edges; ArkAsyncBench targets complementary continuation gaps. Table 6 shows that T1 callback and T5 await paths receive corroborating evidence from both analyses, while all nine T4 Promise-then positives are supplement-exclusive. Removing supplementation therefore preserves 15 core-supported positives but misses the nine T4 cases ($p = 0.00390625$), with no change among the 24 negatives.

A separate diagnostic over 26 callback-heavy projects confirms all eight selected real-code asynchronous paths: a callback metadata property, a Promise-carried advertising identifier, and six

Table 7. Manual semantic audit of 120 real-project paths.

Audited dimension	All	Privacy data	Framework input	Overall rate
Source identity	120/120	90/90	30/30	100.00%
Sink identity	120/120	90/90	30/30	100.00%
Explicit value dependence	114/120	90/90	24/30	95.00%
Reachability	114/120	90/90	24/30	95.00%
Provenance label	120/120	90/90	30/30	100.00%
Complete path	114/120	90/90	24/30	95.00%

Table 8. Privacy evidence across 1,014 projects.

Evidence	Records	Projects	Project rate	Prod./test/gen. paths
Privacy-API usages	2,336	353	34.81%	–
Detector-local sinks	1,754	224	22.09%	–
Privacy-data paths	243	29	2.86%	231/2/10
Framework-input paths	206	159	15.68%	49/157/0
All configured paths	449	181	17.85%	280/159/10

awaited media-metadata or reverse-geocoding values. This links every implemented benchmark construct to executable project code.

Stratified semantic audit. Table 7 reports the 120-path source/IR review. Source identity, sink identity, and provenance are correct for every path. All 90 sampled privacy-data paths are fully correct, including all 30 supplement-only and 26 dual-provenance paths in the full sample. Overall path precision is 114/120 (95.00%; 95% Wilson interval 89.52–97.69%).

The six rejected paths all belong to the optional framework-input model. Four join a lifecycle-written static URI field to unrelated component fields or constants; two connect a lifecycle argument to a permission Toast or a hard-coded event value. This localization matters operationally: the core privacy-data analysis attains 90/90 in the audit, while the framework-input layer remains separately identifiable by source kind and provenance instead of silently entering the privacy-path claim.

Answer to RQ3. IR recovery supplies the largest statistically supported coverage gain, receiver refinement and lifecycle bounds preserve negative-case discrimination, and asynchronous supplementation uniquely recovers nine Promise-then cases. The real-project audit confirms every sampled privacy-data path and measures 95.00% precision across the broader path universe.

5.5 RQ4: Large-Corpus Characterization

The final analysis covers 31,205 ArkTS/TypeScript files and 233,894 methods. Table 8 separates detector evidence, privacy-data paths, and framework-input paths. All 449 configured-query paths retain source and sink endpoints; source scope further distinguishes production, test, and generated code.

The scope split exposes two distinct corpus phenomena. Privacy-data paths overwhelmingly occur in production source (231/243), while framework-input paths are dominated by generated test abilities (157/206). We therefore use the former for privacy-flow claims and retain the latter as a separately labeled analysis of framework-provided values.




Fig. 5. Corpus evidence concentration, category-sink structure, scale, and path complexity.

Concentration and structure. API evidence is long-tailed: 65.19% of projects have no recognized API, while the top 10, 50, and 120 projects contain 24.49%, 54.92%, and 77.83% of all occurrences. The API Gini coefficient is 0.857; HHI is 0.0104, equivalent to 95.8 equally sized contributing projects. Network connectivity (18.28%), screen identity (17.25%), hardware identity (15.11%), account data (8.05%), and clipboard data (5.09%) form the five largest categories.

Project size helps explain identity volume but not path volume. Method count correlates moderately with API occurrences ($\rho = 0.485$), whereas API occurrences correlate strongly with detector-local sinks ($\rho = 0.789$) and only weakly with configured paths ($\rho = 0.175$). From the smallest to largest method quintile, API-positive rates rise from 11.82% to 78.22% and sink-positive rates from 4.43% to 57.43%, while path-positive rates rise to 24.75%. Figure 5 visualizes these concentration, category, scale, and path-complexity relationships.

Destinations and paths. Logging accounts for 1,441/1,754 detector-local sinks (82.16%), followed by UI display (136; 7.75%), storage (94; 5.36%), data return (57; 3.25%), and network (26; 1.48%). Call chains have median/P95 lengths of 1/4; configured paths have median/P95 lengths of 5/10 and a maximum of 21 statements. Provenance comprises 322 IFDS, 84 supplementary, and 43 dual paths.

Integrity and cost. All 1,014 projects produce complete reports. The audit resolves all 3,316 trace endpoints, validates all 236 detector-to-path links, and finds no invalid source kind or provenance. No log reports IFDS batching, budget exhaustion, callback limits, timeout, or out-of-memory failure. The solver processes 3.16 million IFDS edges in total; per-project median/P95/maximum is 805/11,396/197,173. Deduplication reduces 492 raw outcomes to 449 paths. Median runtime is 20.2 s, P95 is 26.9 s, and the 10,757-method outlier completes in 577.9 s.

Clone sensitivity. The token profile contains 31,998 source files and 21,430 unique fingerprints. It forms nine exact-duplicate and eleven near-duplicate project clusters. Retaining one representative per exact cluster changes API/sink/path-positive rates from 34.81/22.09/17.85% to 35.18/22.31/17.99%; near-cluster representatives yield 35.05/22.26/18.03%. The maximum shift is 0.36 percentage points, showing that project-level rates are not driven by repeated copies of a small template set.

Answer to RQ4. The same frozen analysis completes all 1,014 projects without resource fallbacks. Privacy-data paths are primarily production-source phenomena, identity and sink evidence are strongly concentrated, and clone-aware resampling leaves all project rates stable within 0.36 percentage points.

6 Discussion

6.1 Evidence as the Analysis Contract

ARKPRISM treats source resolution as an evidence-producing analysis rather than a configuration side effect. Every accepted API identity carries a package/receiver/access witness; every source fact carries a typed value; and every path records its producing analysis. This contract separates three

questions that privacy tools often collapse: whether a configured API occurs, whether a modeled value may reach a configured sink, and whether that behavior violates a policy. ARKPRISM answers the first two and supplies the source, sink class, permission, call context, and provenance needed to review the third.

The empirical results follow the same structure. ArkSourceFirst60 and ArkIdentityBench evaluate exact-location identity decisions. HapBench evaluates case and endpoint-pair flow classification. ArkAsyncBench isolates continuation recovery, while the 120-path source audit tests real-project path semantics. The Top-120 set stresses recovery of dense, previously confirmed identity evidence. The corpus then measures prevalence, concentration, traceability, and cost without treating tool output as a corpus-wide oracle.

6.2 Core Privacy Paths and Framework Context

The semantic audit provides a useful boundary. All 90 sampled privacy-data paths are correct, including IFDS, supplementary, and dual-provenance results. The six rejected paths occur only when the optional framework-input model composes lifecycle values with component or event context: four confuse unrelated static/component fields, and two retain control co-occurrence without explicit value dependence. Because source kind is part of every record, these candidates remain outside the core privacy-data result and can be filtered without changing API identity or privacy-carrier propagation.

This boundary also clarifies corpus interpretation. Of 243 privacy-data paths, 231 originate in production source; by contrast, 157/206 framework-input paths originate in generated test abilities. Reporting the two source universes separately exposes this structural difference and gives auditors a direct way to prioritize production privacy evidence.

6.3 Precision-Oriented Reachability

Lifecycle modeling illustrates why the evaluation emphasizes specificity and balanced accuracy in addition to recall and F1. Unbounded cycling accepts more cross-event positives under the published HapBench labels, but also creates the suite's explicit unreachable-flow false positive. F1 ignores the lost true negative; balanced accuracy accounts for both classes and favors the bounded policy. Under the executed-order sensitivity contract, the bounded policy classifies all 67 cases, whereas unbounded cycling creates three false positives. The default therefore requires an explicit parent, same-instance, or later-creation witness before connecting lifecycle phases.

Asynchronous supplementation follows the same principle. It is bounded independently from IFDS and does not erase its origin. T1 callback and T5 await paths can be corroborated by both analyses; T4 Promise-then supplies the benchmark's nine unique increments. Provenance makes this composition inspectable rather than presenting every recovered path as an equivalent solver proof.

6.4 Validity Controls

Independent labels. ArkSourceFirst60 is fixed by project-scale and configured-import strata before detector execution. Its project-wide candidate enumeration locates code but does not assign labels, and exact source hashes bind all 186 decisions. ArkAsyncBench similarly defines positives and matched negatives from source-level value dependence. The output-ranked Top-120 set has a narrower role: the final configuration is evaluated on reproduction of 666 confirmed keys, while 182 newly reported keys remain outside its metric.

Configuration and artifact identity. Detector rules, typed sources, framework inputs, sinks, package migration, SDK contents, analyzer source, compiled build, and runner are independently

hashed. HapFlow and ARKPRISM share the artifact’s API-17 SDK and source–sink query on HapBench; real-project experiments use the fixed API-20 SDK with fallback disabled. Evaluators verify report inventories, source/report hashes, endpoint links, source kinds, and provenance before aggregating results.

Corpus composition. The corpus includes applications, libraries, demonstrations, tests, and a small amount of generated source. Results therefore characterize the analyzed collection. Source-scope stratification prevents generated test abilities from being mistaken for production privacy paths. Exact- and near-clone resampling changes project rates by at most 0.36 percentage points, showing that the principal corpus patterns are not an artifact of repeated project copies.

6.5 Applicability

ARKPRISM is a source-based static analysis for ArkTS projects. Native C/C++ bridges, runtime-loaded code, server-side behavior, and consent state require complementary evidence such as HarmoBridge [18] or runtime monitoring. Within its stated contract, ARKPRISM produces reviewable API identities and may-paths rather than automatic legal conclusions.

7 Related Work

7.1 Source Identity and Privacy API Recognition

Accurate source/sink specifications are critical for mobile privacy analysis. Android studies have mined permission mappings and source/sink lists via Stowaway, PScout, SuSi, Axplorer, and API-documentation analyses [6, 7, 21, 43, 45], and audited privacy behavior by connecting code, policies, and permissions [3, 55, 62, 63]. These works show that source/sink definitions encode platform semantics and strongly affect reported results.

ARKPRISM applies this principle to OpenHarmony, where package migration from @ohos.* to consolidated @kit.* exports compounds the challenge. Its matching function separates API identity from exact package spelling while retaining namespace constraints. HapFlow’s released Json2ArkMethodSignature loader resolves namespace/method entries to SDK declarations, but its rule schema does not encode executable fields or detector-side factory-origin witnesses [15]. This is a limitation of the released configuration path, not of method signatures in general: a signature analysis with precise receiver types can represent manager methods. ArkAnalyzer provides ArkTS code representation and call-graph construction [14] but does not define privacy API identity resolution. ARKPRISM adds that evidence-carrying layer.

7.2 Asynchronous Semantics and Lifecycle Modeling

Mobile analysis must model lifecycle, callbacks, and implicit edges precisely. Android ICC modeling includes ComDroid, CHEX, Epicc, IC3, IccTA, DidFail, and RAICC [19, 29, 31, 36, 38, 39, 48]; EdgeMiner recovers implicit framework callbacks [13]. Hybrid-app analyses bridge Java, JavaScript, and WebView semantics [11, 30, 57, 58]. These studies show that missing implicit edges dominate mobile analysis errors.

OpenHarmony has analogous but different callback problems: ArkUI declarative builders, class-field callbacks, FunctionType/ClosureType values, and Promise handlers create edges absent from ordinary calls. WEMINT [54] demonstrates explicit callback modeling for asynchronous mini-program APIs via taint rules; ARKPRISM differs by resolving callbacks from ArkAnalyzer’s FunctionType/ClosureType IR and by handling ArkUI builder-option and receiver-factory semantics. HapFlow handles closure semantics [15]; ARKPRISM adds a bounded, separately provenance-labeled post-IFDS recovery for unresolved callback, Promise .then(), and await paths, plus report-CG

edges for field and builder callbacks. The lifecycle state machine (Section 3.6) captures ability/component phase transitions that the flat DummyMain cannot express.

7.3 Evidence Representation and Large-scale Auditing

Android taint analysis has a rich history. TaintDroid [20] motivated smartphone privacy monitoring; FlowDroid [5] introduced context-/flow-/field-/object-sensitive lifecycle-aware analysis; Amandroid, DroidSafe, IccTA, LeakMiner, and AndroidLeaks explored inter-procedural data-flow, vetting, and market-side detection [25, 26, 31, 52, 61]; HornDroid and COVERT added formal and compositional perspectives [8, 12]. More recently, PacDroid unifies pointer analysis, inter-procedural propagation, ICC, and lifecycle handling and evaluates the resulting soundness–precision–performance tradeoff against several Android frameworks [16]. These systems are important methodological references, but they consume Android bytecode and Android framework models; running them on ArkTS would change both the input language and the analysis query. We therefore do not present their published numbers as an empirical baseline for ARKPRISM.

ARKPRISM uses the author-released HapFlow artifact as the directly comparable OpenHarmony taint baseline because it accepts ArkTS projects and answers the same source-to-sink reachability query. ArkAnalyzer and its context-sensitive Pointer Analysis Kit (APAK) provide the source IR, call graphs, and points-to information on which a taint client can be built [14, 60], but do not themselves emit privacy source-to-sink findings. HarmonyFlow instead analyzes compiled Ark Panda IR and evaluates pointer/call-edge recovery [51]; its input representation and output query make it a complementary front-end reference rather than an end-to-end baseline for this experiment. TaintMini’s Universal Data Flow Graph [33] is a core analysis representation that drives cross-page data-flow resolution in mini-programs; ARKPRISM’s evidence subgraph is an audit-oriented representation assembled from detection, call-chain, sink, and taint results with provenance tracking.

Malware and behavior research uses permissions, API calls, graphs, and models to classify malicious apps [1, 4, 22, 24, 40, 44, 49, 53, 56, 59, 64–66]; AndroZoo enables ecosystem-scale studies [2]. ARKPRISM follows this empirical principle: the information-flow subgraph is more structured than a flat feature vector because auditors need source snippets, call contexts, sink classes, and metadata.

ARKPRISM builds on IFDS/IDE foundations [46, 47] and reusable infrastructures (Soot, WALA, Heros [9, 23, 50]). ArkAnalyzer plays this role for ArkTS [14]. APAK shows that ArkTS closure and framework modeling materially affects downstream call targets and false positives relative to CHA/RTA [60]; ARKPRISM consumes this pointer evidence in receiver-sensitive source and sink resolution. Other OpenHarmony efforts explore testing [17, 35], repair [34], and compiled-IR analysis [51], showing that OpenHarmony requires ecosystem-specific analysis [27, 32, 41]. ARKPRISM follows the same layered principle: reuse the framework, then add domain-specific recognition and provenance-annotated subgraph mapping.

7.4 Evaluation Methodology for Mobile Analyses

Comparative studies show that mobile-analysis results are easily distorted by different source/sink lists, different queries, incomplete benchmark subsets, and undocumented ground truth. ReproDroid evaluates six Android taint analyzers under a common query and execution framework [42]. TaintBench goes further by storing expected and unexpected real-world flows in a machine-readable format and validating labels through multi-expert review [37]. Mutation-based evaluation complements fixed benchmarks by systematically exposing unsound analysis choices [10]. These studies shape our protocol: we use the complete author-released HapBench suite, publish a case-level oracle, keep the source and sink query fixed, count incomplete executions separately, report

both positive- and negative-case metrics, and retain every misclassified case. The high-output real-project audit is reported as a precision-oriented audit rather than an independent recall benchmark because projects were selected using ARKPRISM’s output.

8 Conclusion

ARKPRISM makes privacy-source identity a first-class, auditable stage of ArkTS static analysis. Package migration, receiver origin and type, executable properties, compound namespaces, and IR aliases establish which platform API occurs; typed carriers and provenance preserve how its value enters IFDS or bounded asynchronous recovery. This evidence contract resolves all 90 exact-location occurrences in the source-first benchmark, classifies all 64 matched identity cases, and combines 100% case-level precision and specificity with 97.09% F1 on HapBench. ArkAsyncBench isolates nine supplement-exclusive Promise-then paths, while the real-project audit confirms all 90 sampled privacy-data paths and 95% of the broader path sample. The same frozen analysis completes 1,014 projects, identifies 2,336 privacy-API usages and 243 privacy-data paths, and retains source, sink, and provenance evidence for each result. Together, these results establish evidence-carrying source resolution as a practical foundation for precise ArkTS privacy auditing.

Data Availability

The artifact contains the ARKPRISM source and build scripts; fixed API, source, sink, and package-migration rules; ArkSourceFirst60 and ArkIdentityBench; the Top-120 stress audit; all HapBench and ArkAsyncBench oracles and raw outputs; the 120-path semantic review queue and decisions; and per-project reports for the 1,014-project corpus. Evaluators regenerate every table and figure from retained manifests and machine-readable results. Dataset redistribution follows each constituent project’s license; project identifiers, hashes, and derived reports are supplied where source redistribution is restricted.

A Reproduction Details

A.1 HapBench Case-level Outcomes

Table 9 lists every error produced by either executable tool under the fixed 67-case oracle. Cases absent from the table are classified correctly by both tools. We retain the source-annotated oracle even where a more restrictive event or object-lifetime semantics would suggest a different label.

Table 9. Case-level error inventory for the main HapBench comparison.

Tool	Outcome	Case	Observed mechanism
HapFlow	FN	AnonymousClass2	Anonymous instance initializer omitted
HapFlow	FN	Exceptions4	Exceptional successor/handler not recovered
HapFlow	FP	ArrayIndexNoLeak	Distinct constant indices collapse
HapFlow	FP	VirtualDispatch2	Infeasible receiver target retained
HapFlow	FP	VirtualDispatch3	Infeasible receiver target retained
HapFlow	FP	UnreachableFlow	Lifecycle cycle creates reachability
ARKPRISM	FN	ActivityLifecycle4	Parent source method is bypassed by child override
ARKPRISM	FN	BackupExtensionAbility	Oracle assumes state across extension instances
ARKPRISM	FN	Button1	Oracle assumes click-written static state reaches later creation

The two tools agree on 58 outcomes. Of the nine discordant cases, ARKPRISM alone is correct on six and HapFlow alone on three. The exact two-sided McNemar test is therefore based on the discordant pair (6, 3) and yields $p = 0.508$.

At the finer endpoint-pair unit, ARKPRISM’s 52 raw paths canonicalize to 51 source–sink pairs. Source annotations or direct source/IR review confirm 50; the additional pair is the NoLeak.log target at line 14 of VirtualDispatch1, whose allocation receives an empty constant. The retained audit records every pair, its source identity, endpoint lines, provenance, review mode, and decision; an integrity test reproduces TP/FP/FN= 50/1/3 and 98.04% precision.

A.2 Lifecycle-Oracle Sensitivity

The primary evaluation preserves all published/source-annotated HapBench labels. We additionally define a strict executed-order contract for the three disputed lifecycle cases: an override does not execute its parent body without an explicit super call; an instance field does not cross framework-created instances without a same-object witness; and a UI callback occurring after initial ability creation does not feed an earlier onCreate without an explicit later-creation edge. This contract relabels ActivityLifecycle4, BackupExtensionAbility, and Button1 as negative and changes no other case.

Table 10. Sensitivity to strict lifecycle execution order.

Tool	TP	TN	FP	FN	Precision	Recall	Specificity	F1	MCC
HapFlow [15]	48	10	7	2	87.27%	96.00%	58.82%	91.43%	0.622
ARKPRISM	50	17	0	0	100.00%	100.00%	100.00%	100.00%	1.000

The tools then disagree on nine cases, all classified correctly only by ARKPRISM; the exact two-sided McNemar test gives $p = 0.00390625$. This secondary sensitivity contract does not replace the published oracle; it makes the object and event-order semantics responsible for the difference explicit. The artifact derives both tables from the same case predictions through analyze_hapbench_oracle_sensitivity.js.

A.3 Uncertainty and Category Robustness

Table 11 reports Wilson intervals for the main configurations. The wide specificity intervals are a direct consequence of having only 14 negative cases; balanced accuracy and MCC are included to make this class imbalance visible.

Table 11. HapBench uncertainty and category robustness.

Tool	Precision 95% CI	Recall 95% CI	Specificity 95% CI	Macro F1	Worst-category F1
HapFlow [15]	[82.7, 97.1]	[87.2, 99.0]	[45.4, 88.3]	95.73%	88.89%
ARKPRISM	[92.9, 100.0]	[84.6, 98.1]	[78.5, 100.0]	97.96%	85.71%
–IR recovery	[90.6, 100.0]	[56.5, 80.5]	[78.5, 100.0]	87.07%	66.67%
–receiver refinement	[89.7, 99.7]	[84.6, 98.1]	[68.5, 98.7]	97.57%	85.71%
Unbounded lifecycle	[90.1, 99.7]	[90.1, 99.7]	[68.5, 98.7]	98.81%	91.67%

Table 12. Per-category F1 on all HapBench cases.

Category	Cases	HapFlow	ARKPRISM
Aliasing	6	100.00%	100.00%
Anonymous constructs	10	93.33%	100.00%
Array-like structures	5	88.89%	100.00%
Field/object sensitivity	6	100.00%	100.00%
General language features	22	91.89%	100.00%
Lifecycle modeling	13	96.00%	85.71%
OpenHarmony-specific APIs	5	100.00%	100.00%

A.4 Source-first Benchmark Protocol

ArkSourceFirst60 is selected before the evaluated detector run. Projects are partitioned by source-tree size and by whether their imports are compatible with the fixed sensitive-API rules; deterministic hash ordering selects ten projects from each of the six strata. The manifest stores the selection seed, source-tree fingerprint, rule hash, and package-migration hash. Candidate discovery collects the project-wide SDK-import universe and enumerates statically named calls and property reads compatible with that universe, including static element access, named-import calls, and simple callable aliases. It neither reads a ARKPRISM report nor assigns a gold label.

Manual review accepts an occurrence only when its executable source context establishes the configured package/namespace/member and access kind through an import, SDK receiver type, constructor, manager/factory definition, or executable property read. It rejects comments, strings, declarations, same-name application methods, and unrelated library/controller receivers. The delivered gold file contains 90 accepted occurrences; a companion decision file preserves all 186 judgments, including 96 receiver- or owner-incompatible negatives. An integrity test rehashes every source file, validates every exact line and column, checks each canonical rule key, and rejects duplicate occurrence identities. Evaluation joins detector and gold records on project, normalized file, exact line and column, canonical API, and access kind.

A.5 High-output Stress-audit Protocol

The 120-project stress audit labels project–namespace–member keys. A key is accepted only when executable source establishes the configured API through an imported namespace/package alias, a

manager/helper definition chain, an SDK receiver type, or an executable property rule. Comments, strings, declarations, unrelated same-name methods, and ambiguous method-only matches are rejected. Each accepted key retains project, normalized package/namespace/member, access kind, source file/line, and a bounded source excerpt.

The construction set contains 1,533 reviewed location rows that canonicalize to 666 confirmed project-API keys. The companion integrity audit resolves every key against the delivered source tree and final rules, verifies file/line bounds and snippets, and checks package/namespace/member compatibility. All 666 records pass. Because the set begins from an earlier output ranking, it is not an exhaustive occurrence oracle.

The final detector-only run uses the same frozen build as the corpus and recovers all 666 confirmed keys in all 120 projects. It also reports 182 keys outside the frozen review set. The evaluator therefore defines *confirmed-key reproduction coverage* (666/666) and leaves final-output precision undefined; no unreviewed key enters either side of a precision ratio.

A.6 Real-project Path-audit Protocol

The path audit begins from the complete 449-path corpus artifact. A deterministic sampler balances source kind, provenance, sink family, path-length bin, and project diversity, producing 120 records from 48 projects. The fixed queue contains 90 privacy-data and 30 framework-input paths; 64 IFDS, 30 supplementary, and 26 dual-provenance paths; and all available network and UI strata. Each record stores the run-manifest hash, report hash, source/sink file hashes, endpoint locations, path statements, source identity, and provenance.

Reviewers inspect the referenced source and IR path and assign five Boolean decisions: source identity, sink identity, explicit value dependence, reachability consistency, and provenance label. The evaluator accepts a complete path only when all five hold, rejects missing decisions, and rehashes all 240 endpoint files and 120 reports before computing metrics. The audit confirms source identity, sink identity, and provenance for 120/120 paths; explicit dependence, reachability, and complete-path correctness hold for 114/120. All 90 privacy-data paths are complete. The six rejected framework-input paths consist of four unrelated static/component-field joins and two lifecycle/event joins without explicit value dependence.

A.7 Artifact-to-Paper Traceability

The artifact provides ArkSourceFirst60's selection manifest, complete 186-item decision file, occurrence gold set, candidate-enumerator tests, and integrity evaluator; the 64-case ArkIdentityBench oracle; machine-readable oracles and raw outputs for all 67 HapBench and 48 ArkAsyncBench cases; the Top-120 reproduction evaluator; and the complete semantic-path audit. Evaluators derive confusion matrices, construct/category metrics, endpoint-pair outcomes, paired tests, uncertainty intervals, and runtimes. Figure 4 and Figure 5 are generated from evaluator JSON rather than transcribed values.

Corpus tables and figures consume one verified report per project from a single frozen build, SDK, rule set, and command contract; change-impact overlays are rejected by the release summarizer. The manifest records project inventory, build/source/runner hashes, SDK and rule fingerprints, resource bounds, and completion state. The corpus audit checks that every configured path has resolvable source and sink endpoints, a valid ifds, async_supplement, or both provenance tag, and valid detector-to-path links when the two evidence universes can be joined. Companion JSON retains every plotted value and input manifest hash.

References

- [1] Yousra Aafer, Wenliang Du, and Heng Yin. 2013. DroidAPIMiner: Mining API-Level Features for Robust Malware Detection in Android. In *Proceedings of the International Conference on Security and Privacy in Communication Systems*. 86–103.
- [2] Kevin Allix, Tegawendé F. Bissyandé, Jacques Klein, and Yves Le Traon. 2016. AndroZoo: Collecting Millions of Android Apps for the Research Community. In *Proceedings of the 13th International Conference on Mining Software Repositories (ICMSR)*. ACM, New York, NY, USA, 468–471.
- [3] Benjamin Andow, Samin Yaseer Mahmud, Justin Whitaker, William Enck, Bradley Reaves, Kapil Singh, and Serge Egelman. 2020. Actions Speak Louder than Words: Entity-Sensitive Privacy Policy and Data Flow Analysis with PoliCheck. In *29th USENIX Security Symposium (USENIX Security)*. USENIX Association, Online, 985–1002.
- [4] Daniel Arp, Michael Spreitzenbarth, Malte Hubner, Hugo Gascon, Konrad Rieck, and CERT Siemens. 2014. DREBIN: Effective and Explainable Detection of Android Malware in Your Pocket. In *Proceedings of the Network and Distributed System Security Symposium*.
- [5] Steven Arzt, Siegfried Rasthofer, Christian Fritz, Eric Bodden, Alexandre Bartel, Jacques Klein, Yves Le Traon, Damien Ochteau, and Patrick McDaniel. 2014. FlowDroid: precise context, flow, field, object-sensitive and lifecycle-aware taint analysis for Android apps. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. Association for Computing Machinery, New York, NY, USA, 259–269.
- [6] Kathy Wain Yee Au, Yi Fan Zhou, Zhen Huang, and David Lie. 2012. PScout: Analyzing the Android Permission Specification. In *Proceedings of the ACM Conference on Computer and Communications Security*. 217–228.
- [7] Michael Backes, Sven Bugiel, Erik Derr, Patrick McDaniel, Damien Ochteau, and Sebastian Weisgerber. 2016. On Demystifying the Android Application Framework: Re-Visiting Android Permission Specification Analysis. In *Proceedings of the USENIX Security Symposium*. 1101–1118.
- [8] Hamid Bagheri, Alireza Sadeghi, Joshua Garcia, and Sam Malek. 2015. COVERT: Compositional Analysis of Android Inter-App Permission Leakage. In *Proceedings of the IEEE/ACM International Conference on Software Engineering*. 866–877.
- [9] Eric Bodden. 2012. Inter-procedural Data-flow Analysis with IFDS/IDE and Soot. In *Proceedings of the ACM SIGPLAN International Workshop on State of the Art in Java Program Analysis*. 3–8.
- [10] Richard Bonett, Kaushal Kafle, Kevin Moran, Adwait Nadkarni, and Denys Poshyvanyk. 2018. Discovering Flaws in Security-Focused Static Analysis Tools for Android Using Systematic Mutation. In *Proceedings of the 27th USENIX Security Symposium*. 1263–1280. <https://www.usenix.org/conference/usenixsecurity18/presentation/bonett>
- [11] Achim D. Brucker and Michael Herzberg. 2018. BabelView: Evaluating the Impact of Code Injection Attacks in Mobile Webviews. In *Proceedings of the International Symposium on Research in Attacks, Intrusions, and Defenses*. 25–46.
- [12] Stefano Calzavara, Ilya Grishchenko, and Matteo Maffei. 2016. HornDroid: Practical and Sound Static Analysis of Android Applications by SMT Solving. In *Proceedings of the IEEE European Symposium on Security and Privacy*. 47–62.
- [13] Yinzhi Cao, Yanick Fratantonio, Antonio Bianchi, Antonio Chung, Yanick Huang, David Lie, Giovanni Vigna, and Christopher Kruegel. 2015. EdgeMiner: Automatically Detecting Implicit Control Flow Transitions through the Android Framework. In *Proceedings of the Network and Distributed System Security Symposium*.
- [14] Haonan Chen, Daihang Chen, Yizhuo Yang, Lingyun Xu, Liang Gao, Mingyi Zhou, Chunming Hu, and Li Li. 2025. ArkAnalyzer: The Static Analysis Framework for OpenHarmony. In *Proceedings of the 47th International Conference on Software Engineering: Software Engineering in Practice (ICSE SEIP)*. ACM.
- [15] Haonan Chen, Mingyi Zhou, Yanjie Zhao, and Li Li. 2026. HapFlow: The Taint Analysis Framework for OpenHarmony Apps. *ACM Transactions on Software Engineering and Methodology* (2026). doi:10.1145/3793863
- [16] Menglong Chen, Tian Tan, Minxue Pan, and Yue Li. 2025. PacDroid: A Pointer-Analysis-Centric Framework for Security Vulnerabilities in Android Apps. In *Proceedings of the 47th IEEE/ACM International Conference on Software Engineering*. 2803–2815. doi:10.1109/ICSE55347.2025.00208
- [17] Yige Chen, Yujia Fan, Siyi Wang, Yi Tao, and Yepang Liu. 2025. HmTest: Automated Testing of HarmonyOS Apps via Model-Driven Navigation and Reinforcement Learning. *Journal of Computer Science and Technology* 40, 4 (2025), 1006–1021. doi:10.1007/s11390-025-5142-4
- [18] Zhen Chen, Yi Li, Teng Wang, et al. 2025. HarmoBridge: Static Analysis of Cross-Language Data Flows in HarmonyOS Applications. In *Proceedings of the 40th IEEE/ACM International Conference on Automated Software Engineering (ASE)*.
- [19] Erika Chin, Adrienne Porter Felt, Kate Greenwood, and David Wagner. 2011. Analyzing Inter-Application Communication in Android. In *Proceedings of the 9th International Conference on Mobile Systems, Applications, and Services*. 239–252.
- [20] William Enck, Peter Gilbert, Seungyeop Han, Vasant Tendulkar, Byung-Gon Chun, Landon P. Cox, Jaeyeon Jung, Patrick McDaniel, and Anmol N. Sheth. 2014. TaintDroid: An Information-Flow Tracking System for Realtime Privacy Monitoring on Smartphones. *ACM Trans. Comput. Syst.* 32, 2 (2014), 5:1–5:29.

- [21] Adrienne Porter Felt, Erika Chin, Steve Hanna, Dawn Song, and David Wagner. 2011. Android Permissions Demystified. In *Proceedings of the ACM Conference on Computer and Communications Security*. 627–638.
- [22] Yu Feng, Saswat Anand, Isil Dillig, and Alex Aiken. 2014. Apposcopy: Semantics-Based Detection of Android Malware through Static Analysis. In *Proceedings of the ACM SIGSOFT International Symposium on Foundations of Software Engineering*. 576–587.
- [23] Stephen J. Fink, Julian Dolby, and Eran Yahav. 2008. Effective Typestate Verification in the Presence of Aliasing. *ACM Transactions on Software Engineering and Methodology* 17, 2, Article 9 (2008).
- [24] Yanick Fratantonio, Antonio Bianchi, William Robertson, Dhillung Kirat, Christopher Kruegel, and Giovanni Vigna. 2016. TriggerScope: Towards Detecting Logic Bombs in Android Applications. In *Proceedings of the IEEE Symposium on Security and Privacy*. 377–396.
- [25] Clint Gible, Jonathan Crussell, Jeremy Erickson, and Hao Chen. 2012. AndroidLeaks: Automatically Detecting Potential Privacy Leaks in Android Applications on a Large Scale. In *Proceedings of the 5th International Conference on Trust and Trustworthy Computing*. 291–307.
- [26] Michael I. Gordon, Deokhwan Kim, Jeff H. Perkins, Limei Gilham, Nguyen Nguyen, and Martin Rinard. 2015. Information Flow Analysis of Android Applications in DroidSafe. In *Proceedings of the Network and Distributed System Security Symposium*.
- [27] Huawei. 2021. HarmonyOS 2 Security Technical White Paper. <https://consumer.huawei.com/>
- [28] Huawei Developer. 2026. ArkTS: Dedicated Programming Language for the HarmonyOS Ecosystem. <https://developer.huawei.com/consumer/en/arkts/>
- [29] William Klieber, Lori Flynn, Amar Bhosale, Limin Jia, and Lujo Bauer. 2014. Android Taint Flow Analysis for App Sets. In *Proceedings of the 3rd ACM SIGPLAN International Workshop on the State of the Art in Java Program Analysis*. 1–6.
- [30] Sungho Lee, Julian Dolby, and Sukyoung Ryu. 2016. HybriDroid: static analysis framework for Android hybrid applications. In *Proceedings of the 31st IEEE/ACM International Conference on Automated Software Engineering (ASE)*. Association for Computing Machinery, New York, NY, USA, 250–261. doi:10.1145/2970276.2970368
- [31] Li Li, Alexandre Bartel, Tegawendé F. Bissyandé, Jacques Klein, Yves Le Traon, Steven Arzt, Siegfried Rasthofer, Eric Bodden, Damien Ocateau, and Patrick McDaniel. 2015. IccTA: Detecting Inter-Component Privacy Leaks in Android Apps. In *Proceedings of the 37th International Conference on Software Engineering*. 280–291.
- [32] Li Li, Xiang Gao, Hailong Sun, Chunming Hu, Xiaoyu Sun, Haoyu Wang, Haipeng Cai, Ting Su, Xiapu Luo, Tegawendé F. Bissyandé, Jacques Klein, John Grundy, Tao Xie, Haibo Chen, and Huaimin Wang. 2025. Software Engineering for OpenHarmony: A Research Roadmap. *Comput. Surveys* 58, 2, Article 34 (2025), 34 pages. doi:10.1145/3720538
- [33] Yi Li, Yanyan Chen, Teng Wang, et al. 2023. TaintMini: Detecting Privacy Data Leaks in Mini-Programs via Universal Data Flow Graph. In *Proceedings of the 45th International Conference on Software Engineering (ICSE)*.
- [34] Zhihao Lin, Mingyi Zhou, Wei Ma, Chi Chen, Yun Yang, Jun Wang, Chunming Hu, and Li Li. 2025. HapRepair: Learn to Repair OpenHarmony Apps. In *Proceedings of the 33rd ACM International Conference on the Foundations of Software Engineering, Industry Track*. 319–330. doi:10.1145/3696630.3728556
- [35] Farong Liu, Mingyi Zhou, Yakun Zhang, Ting Su, Bo Sun, Jacques Klein, Xiang Gao, and Li Li. 2025. HapTest: The Dynamic Analysis Framework for OpenHarmony. In *Proceedings of the 33rd ACM International Conference on the Foundations of Software Engineering, Companion Proceedings*. 1–10. doi:10.1145/3696630.3728565
- [36] Long Lu, Zhichun Li, Zhenyu Wu, Wenke Lee, and Guofei Jiang. 2012. CHEX: Statically Vetting Android Apps for Component Hijacking Vulnerabilities. In *Proceedings of the ACM Conference on Computer and Communications Security*. 229–240.
- [37] Linghui Luo, Felix Pauck, Goran Piskachev, Manuel Benz, Ivan Pashchenko, Martin Mory, Eric Bodden, Ben Hermann, and Fabio Massacci. 2022. TaintBench: Automatic Real-World Malware Benchmarking of Android Taint Analyses. *Empirical Software Engineering* 27, 1 (2022), 16. doi:10.1007/s10664-021-10013-5
- [38] Damien Ocateau, Somesh Jha, and Patrick McDaniel. 2013. Effective Inter-Component Communication Mapping in Android with Epicc: An Essential Step Towards Holistic Security Analysis. In *Proceedings of the 22nd USENIX Security Symposium*. 543–558.
- [39] Damien Ocateau, Daniel Luchaup, Matthew Dering, Somesh Jha, and Patrick McDaniel. 2015. Composite Constant Propagation: Application to Android Inter-Component Communication Analysis. In *Proceedings of the 37th International Conference on Software Engineering*. 77–88.
- [40] Lucky Onwuzurike, Enrico Mariconti, Panagiotis Andriotis, Emiliano De Cristofaro, Gordon Ross, and Gianluca Stringhini. 2019. MaMaDroid: Detecting Android Malware by Building Markov Chains of Behavioral Models (Extended Version). *ACM Transactions on Privacy and Security (TOPS)* 22, 2, Article 10 (2019), 10:1–10:29 pages.
- [41] OpenAtom Foundation. 2026. OpenHarmony: A Comprehensive Open Source Project for the All-scenario, Fully-connected, Intelligent Era. <https://www.openharmony.cn/>
- [42] Felix Pauck, Eric Bodden, and Heike Wehrheim. 2018. Do Android Taint Analysis Tools Keep Their Promises?. In *Proceedings of the 26th ACM Joint European Software Engineering Conference and Symposium on the Foundations of*

- Software Engineering*. 331–341. doi:10.1145/3236024.3236029
- [43] Siegfried Rasthofer, Steven Arzt, and Eric Bodden. 2014. A Machine-learning Approach for Classifying and Categorizing Android Sources and Sinks. In *Proceedings of the Network and Distributed System Security Symposium*.
- [44] Siegfried Rasthofer, Steven Arzt, Enrico Lovat, and Eric Bodden. 2016. HARVESTER: Automatic Generation of Entry Points for Dynamic Analysis of Android Applications. In *Proceedings of the Network and Distributed System Security Symposium*.
- [45] Talia Ravitch, Eric Creswick, Aaron Tomb, Adam Foltzer, Trevor Elliott, and Ledah Casburn. 2014. Multi-App Security Analysis with Fuse: Statically Detecting Android App Collusion. In *Proceedings of the ACM SIGPLAN International Workshop on the State of the Art in Java Program Analysis*. 4–9.
- [46] Thomas Reps, Susan Horwitz, and Mooly Sagiv. 1995. Precise Interprocedural Dataflow Analysis via Graph Reachability. In *Proceedings of the 22nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. 49–61.
- [47] Mooly Sagiv, Thomas Reps, and Susan Horwitz. 1996. Precise Interprocedural Dataflow Analysis with Applications to Constant Propagation. *Theoretical Computer Science* 167, 1–2 (1996), 131–170.
- [48] Jordan Samhi, Alexandre Bartel, Tegawendé F. Bissyandé, and Jacques Klein. 2021. RAICC: Revealing Atypical Inter-Component Communication in Android Apps. *Empirical Software Engineering* 26, 6 (2021).
- [49] Kimberly Tam, Ali Feizollah, Nor Badrul Anuar, Rosli Salleh, and Lorenzo Cavallaro. 2015. CopperDroid: Automatic Reconstruction of Android Malware Behaviors. In *Proceedings of the Network and Distributed System Security Symposium*.
- [50] Raja Vallée-Rai, Phong Co, Etienne Gagnon, Laurie Hendren, Patrick Lam, and Vijay Sundaresan. 2010. Soot: A Java Bytecode Optimization Framework. In *CASCON First Decade High Impact Papers*. IBM Corp., 214–224.
- [51] Yue Wang, Sen Chen, Jiaming Li, Wenjing Dang, Renchao Chai, Ziyue Shi, and Li Li. 2025. HarmonyFlow: A Static Analysis Framework for HarmonyOS Applications Based on Ark Panda IR. *Journal of Software* (2025), 1–19. doi:10.13328/j.cnki.jos.007477
- [52] Fengguo Wei, Sankardas Roy, Xinming Ou, and Robby. 2018. Amandroid: A Precise and General Inter-component Data Flow Analysis Framework for Security Vetting of Android Apps. *ACM Trans. Priv. Secur.* 21, 3 (2018), 14:1–14:32.
- [53] Michelle Y. Wong and David Lie. 2016. IntelliDroid: A Targeted Input Generator for the Dynamic Analysis of Android Malware. In *Proceedings of the Network and Distributed System Security Symposium*.
- [54] Haoyu Wu, Lingxiang Yang, Mengfei Wang, et al. 2023. WEMINT: Detecting Sensitive Data Flows in Mini-Programs via Explicit Asynchronous Callback Modeling. In *Proceedings of the 38th IEEE/ACM International Conference on Automated Software Engineering (ASE)*.
- [55] Meng Xia, Lu Gong, Yuan Lyu, Zhengwei Qi, and Xue Liu. 2015. Effective Real-time Android Application Auditing. In *Proceedings of the IEEE Symposium on Security and Privacy*. 899–914.
- [56] Lok Kwong Yan and Heng Yin. 2012. DroidScope: seamlessly reconstructing the OS and Dalvik semantic views for dynamic Android malware analysis. In *Proceedings of the 21st USENIX Conference on Security Symposium (USENIX Security)* (Bellevue, WA). USENIX Association, USA, 29.
- [57] Guangliang Yang, Abner Mendoza, and Guofei Gu. 2018. Automated Generation of Event-Oriented Exploits in Android Hybrid Apps. In *Proceedings of the Network and Distributed System Security Symposium*.
- [58] Guangliang Yang, Abner Mendoza, Jialong Zhang, and Guofei Gu. 2017. BridgeScope: Precisely and Scalably Vetting JavaScript Bridge Security Issues in Android Hybrid Apps. In *Proceedings of the International Symposium on Research in Attacks, Intrusions, and Defenses*. 166–185.
- [59] Wei Yang, Xusheng Xiao, Benjamin Andow, Sihan Li, Tao Xie, and William Enck. 2015. AppContext: Differentiating Malicious and Benign Mobile App Behaviors Using Context. In *Proceedings of the International Conference on Software Engineering*. 303–313.
- [60] Yizhuo Yang, Lingyun Xu, Mingyi Zhou, and Li Li. 2026. Context-Sensitive Pointer Analysis for ArkTS. arXiv preprint. <https://arxiv.org/abs/2602.00457>
- [61] Zheming Yang and Min Yang. 2012. LeakMiner: Detect Information Leakage on Android with Static Taint Analysis. In *Proceedings of the 3rd World Congress on Software Engineering*. 101–104.
- [62] Zheming Yang, Min Yang, Yuan Zhang, Guofei Gu, Peng Ning, and X. Sean Wang. 2013. AppIntent: Analyzing Sensitive Data Transmission in Android for Privacy Leakage Detection. In *Proceedings of the ACM Conference on Computer and Communications Security*. 1043–1054.
- [63] Le Yu, Xiapu Luo, Jiachi Chen, Hao Zhou, Tao Zhang, Henry Chang, and Hareton K. N. Leung. 2021. PPChecker: Towards Accessing the Trustworthiness of Android Apps’ Privacy Policies. *IEEE Transactions on Software Engineering (TSE)* 47, 2 (2021), 221–242.
- [64] Mu Zhang, Yue Duan, Heng Yin, and Zhiruo Zhao. 2014. Semantics-Aware Android Malware Classification Using Weighted Contextual API Dependency Graphs. In *Proceedings of the ACM Conference on Computer and Communications Security*. 1105–1116.

- [65] Wu Zhou, Yajin Zhou, Xuxian Jiang, and Peng Ning. 2012. Detecting Repackaged Smartphone Applications in Third-party Android Marketplaces. In *Proceedings of the ACM Conference on Data and Application Security and Privacy*. 317–326.
- [66] Yajin Zhou, Zhi Wang, Wu Zhou, and Xuxian Jiang. 2012. DroidRanger: Systematic Characterization and Detection of Malicious Applications in Official and Alternative Android Markets. In *Proceedings of the Network and Distributed System Security Symposium*.